

# ECON 441B : Intro Machine Learning Lab

Week 1, Lecture 1 | Github and Google Colab

Sam Borghese

Wednesday, January 9th, 2024

1. Course Overview
2. Github
3. Google Colab
4. In-Class Assignment

# Course Overview

# Econ 441B : ~~Applied Data Management For Economists~~

## Econ 441B : Introductory Machine Learning Lab

### **Course Description:**

This is an introductory lab to Machine Learning for Economists. The content will be taught through visualization with real data and pragmatic in class assignments with real or fabricated data. The tasks of the assignments will match the most popular fields for Quantitative Economists mainly focussing on quantitative financial analytics and consumer analytics. The main goal of the course is to arm students with a set of critical thinking skills and algorithms to solve any data problem using a machine learning approach.

The course will synchronize with Econ 425 and provide hands-on examples for the concepts learned inside of that course. The implications of the examples in 441B will also be useful in the Asset Pricing course which will take the tools learned in Machine Learning and apply them to Quantitative Finance. While the ML and Asset Pricing programs synergize they can be taken independently.

**This course was introduced to give an idea of how  
the algorithms are used outside of the textbook and academia**

# Econ 441B : Introductory Machine Learning Lab

This course aligns with Professor Grigory's 425 and Professor Guang's 425

All courses are independent, but I assume you attended their weekly lectures before my course

This course counts for 20% of 425 Grade  
441B Grade = 425 Grade

No curve in this class but maybe in 425

# Econ 441B : Introductory Machine Learning Lab

## Topics

You learn the topics in their class. Practice them in my class

### Topics:

- Week 1: Introduction & Machine Learning Pipeline
- Week 2: Discrete Classification I - Logistic Regression
- Week 3: Linear Models, Regularization, and Hyperparameter Tuning
- Week 4: Discrete Classification II - Decision Trees
- Week 5: Dealing with Imbalanced Data
- Week 6: Neural Nets
- Week 7: Large-Language Models
- Week 8: LLM Cont. & Bagging and Boosting (Non-Parametric Models)
- Week 9: Unsupervised Learning - Clustering & PCA
- Week 10: Reinforcement Learning

# Econ 441B : Introductory Machine Learning Lab

## Assignments

### **Computation of Course Grade:**

- 8 - 12.5% In-Class Assignments (There are 10 total, you can miss 2 of them)
- All work must be posted to GitHub to count
- This grade will count towards 20% of 425 grade
- The grade you get in 425 will be your 441B grade as well

There will be no curve. In-class assignments graded for correctness, with partial credit given.

Each class you will be given a short challenge problem that you will have 10-30 minutes to complete

Assignments should be completed in class but you will have until Friday to Submit  
No Late Work Accepted

# Econ 441B : Introductory Machine Learning Lab

## Guest Presentations

Twice there will be presentations from industry people who use Machine Learning in their jobs commonly. They will share what algorithms, softwares and packages they use to solve various problems.

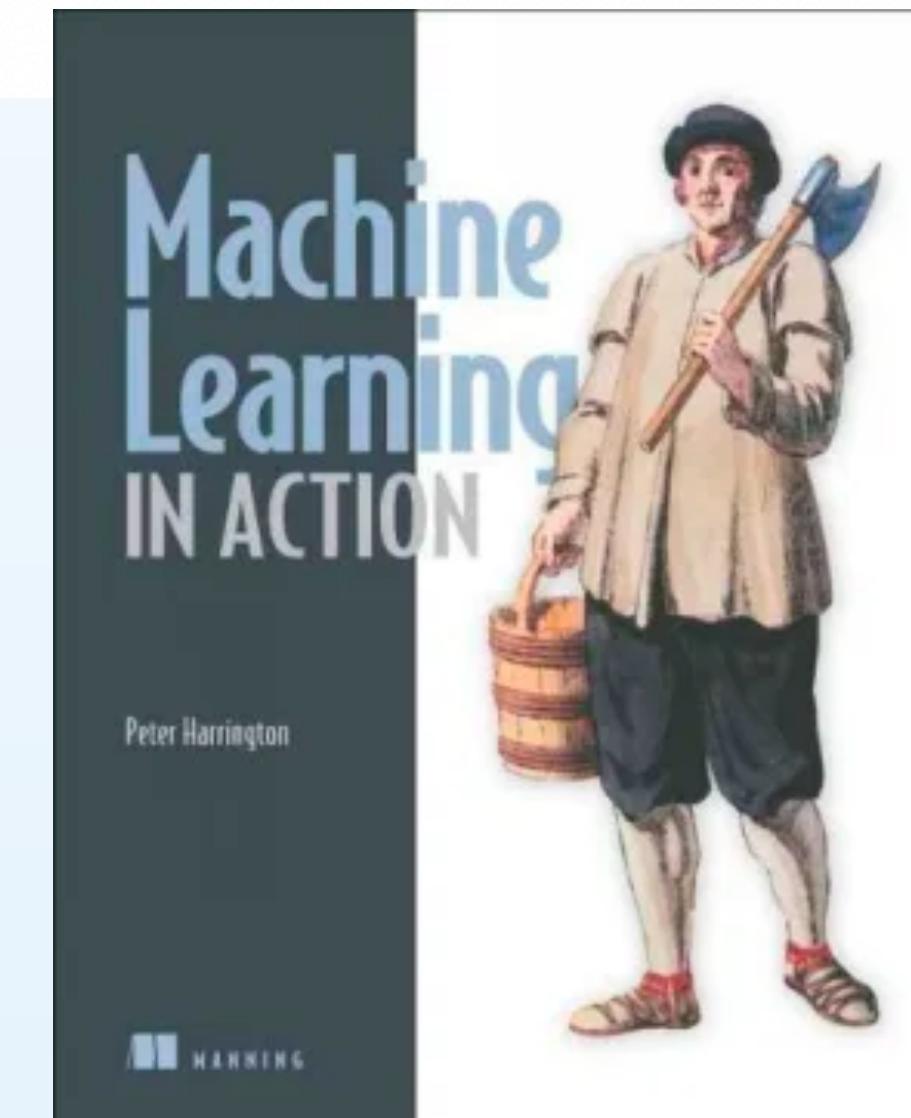
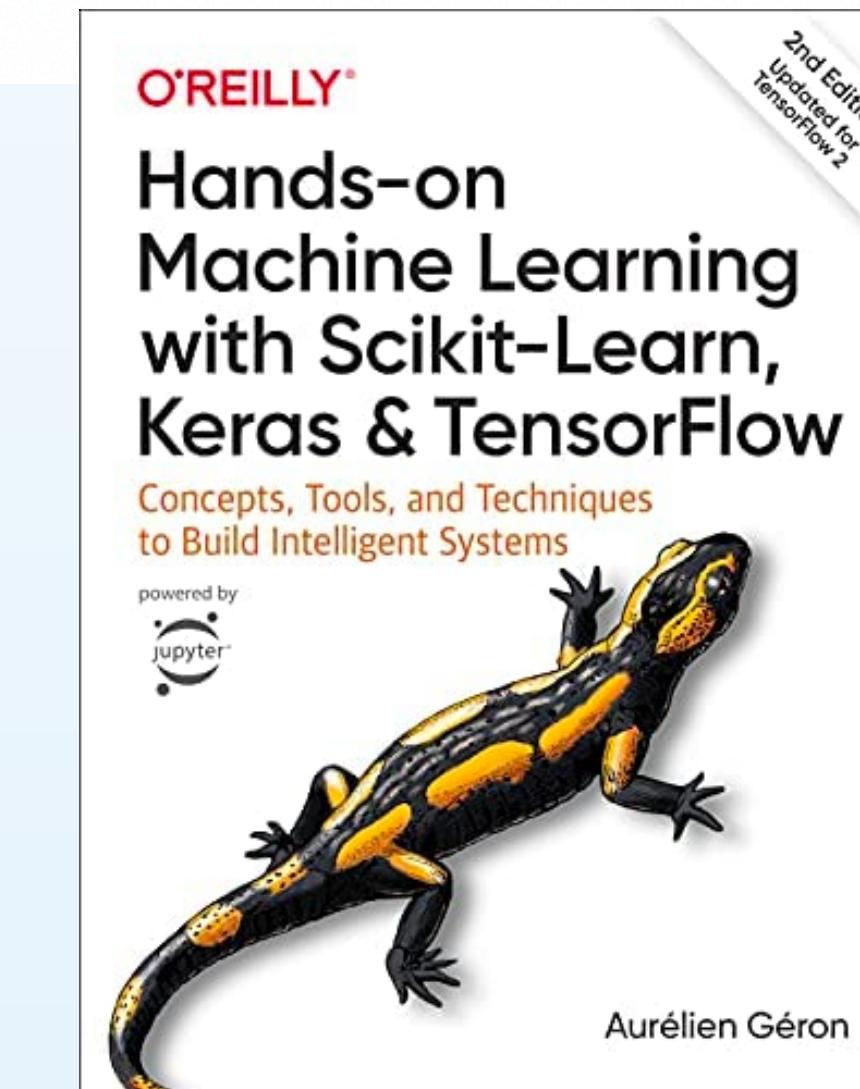
# Econ 441B : Introductory Machine Learning Lab

## Recommended Textbooks

The textbooks are optional. Some of the content of this course is inspired from these books loosely

### Optional Readings:

1. *Hands-On Machine Learning with Scikit-Learn, Keras and TensorFlow*, 2nd Edition (O'Reilly). Copyright: September 2019. [10-day Trial ↗](#) Or free through UCLA Library
2. *Machine Learning in Action* (Peter Harrington). Copyright: April 3, 2012. ISBN : 9781638352457-1638352453



# Github

# What is GitHub?

An online repository that helps developers store, manage, track, share and control their code



## Main benefits :

- 1.) Use other open source code**
- 2.) Share Code with a Team**
- 3.) Version Control**

# How necessary is GitHub?

The most popular source

It is almost a given that you know about GitHub in a Data role



**It is commonly used to host open source software development projects. As of June 2022, GitHub reported having over 83 million developers and more than 200 million repositories, including at least 28 million public repositories.**

**Of the 840 machine learning engineers in the study, 61% said they use GitHub for sharing, the highest percentage of any profession in the report to do so.** Jul 11, 2022

**It is the largest [source code](#) host as of November 2021.**

Source : <https://thenewstack.io/where-do-data-practitioners-prefer-to-collaborate-github/#:~:text=Of%20the%20840%20machine%20learning,the%20report%20to%20do%20so.>

# How necessary is GitHub?

## How Big Tech Contributes to Open Source

Active GitHub contributors by employer by Dec 31st 2020\*



\* Active contributors: More than ten commits on GitHub

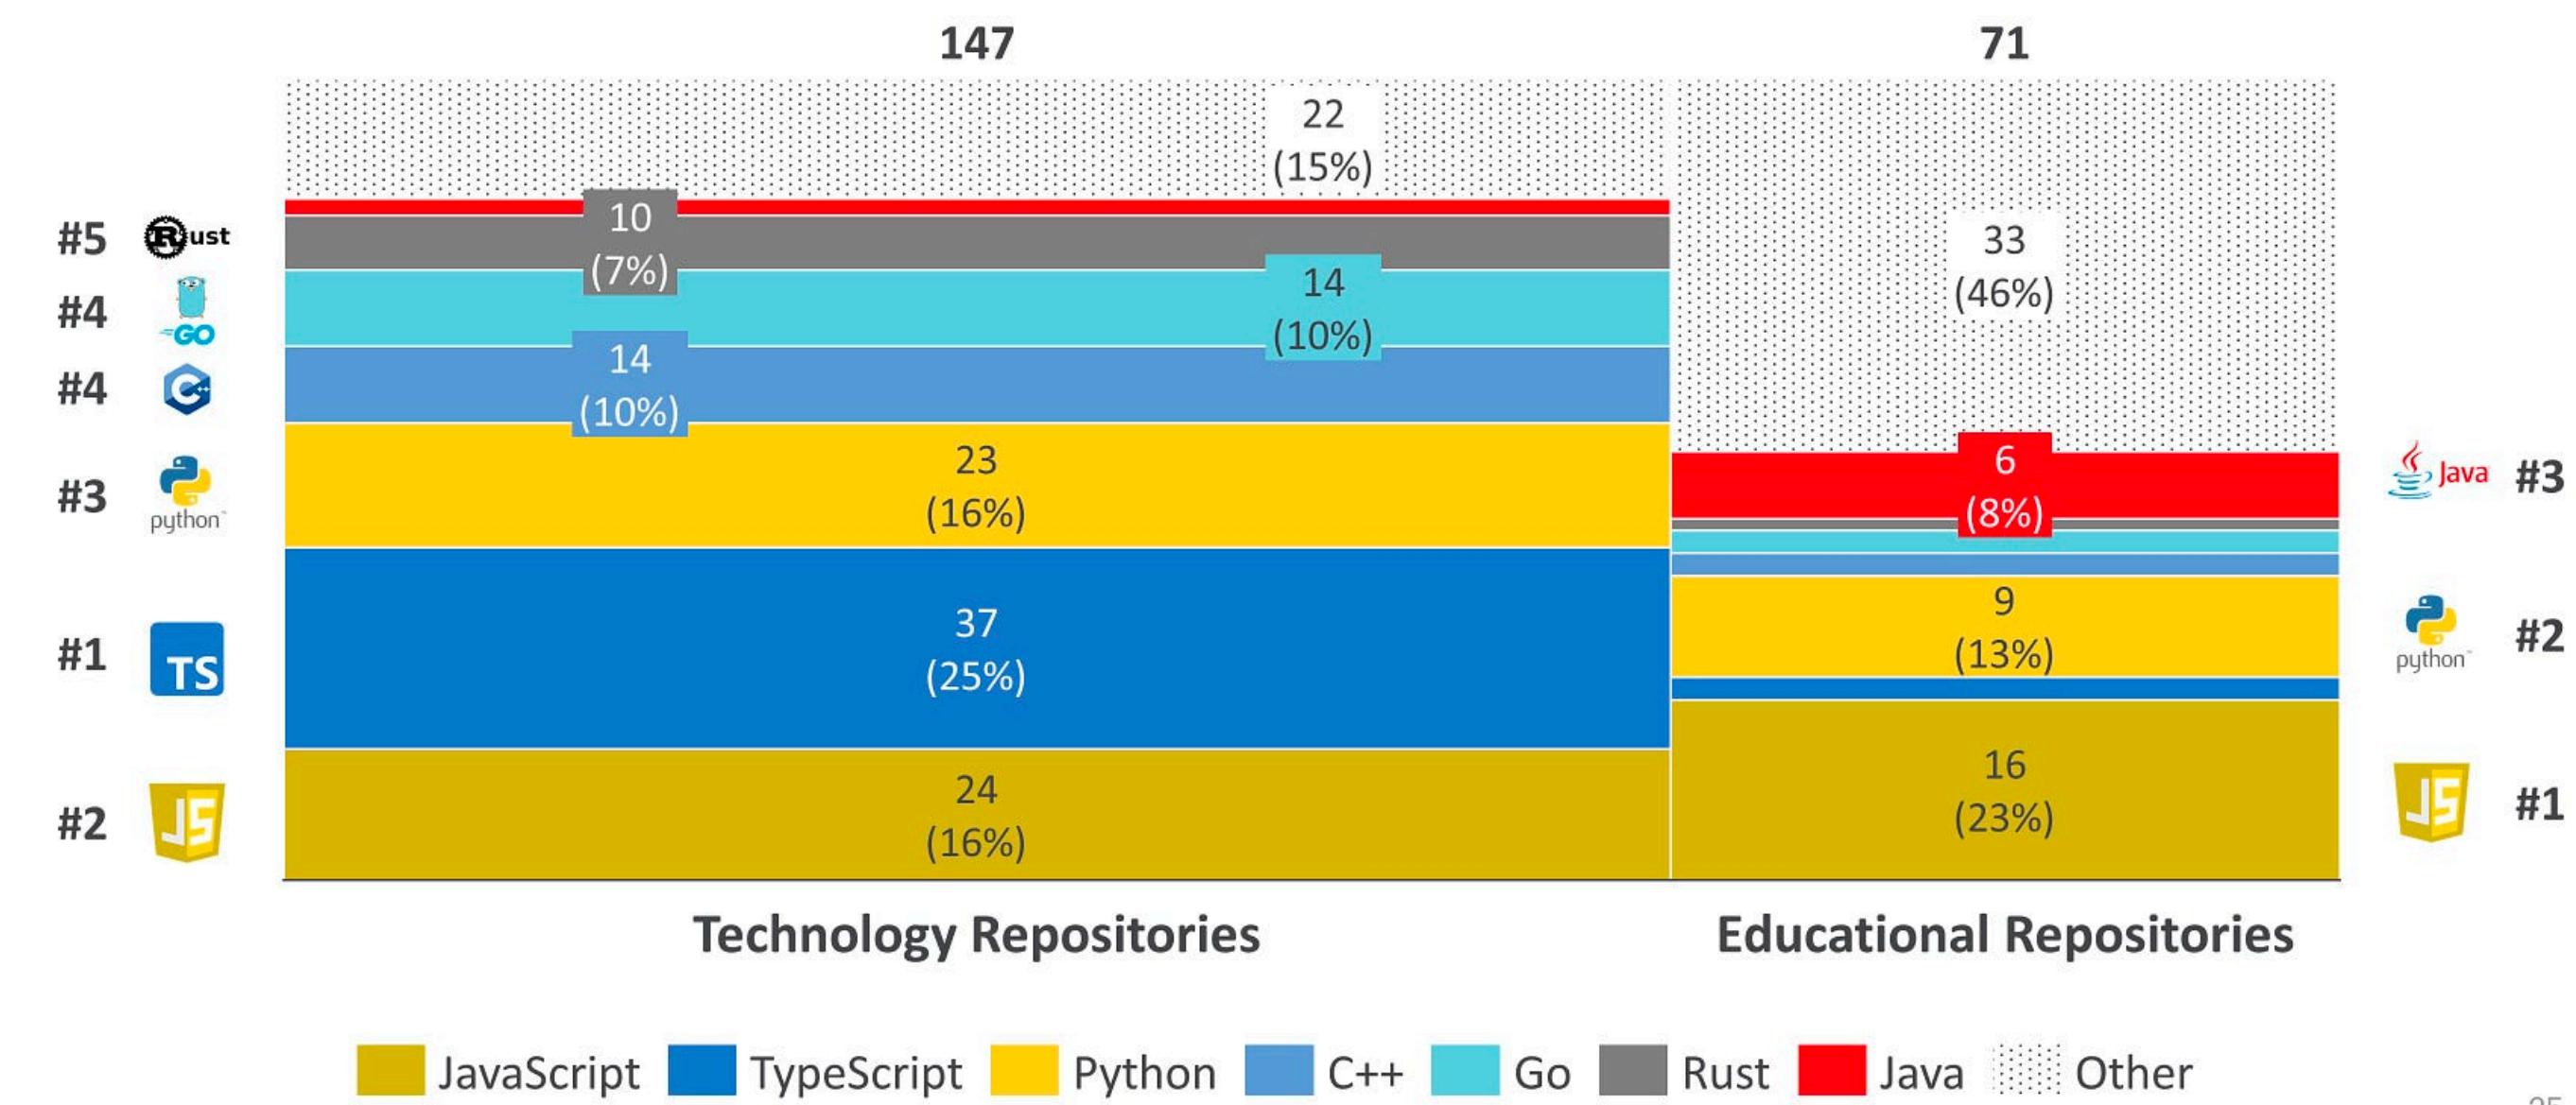
Source: GitHub



statista

## Top programming languages of trending Github repos in 2021

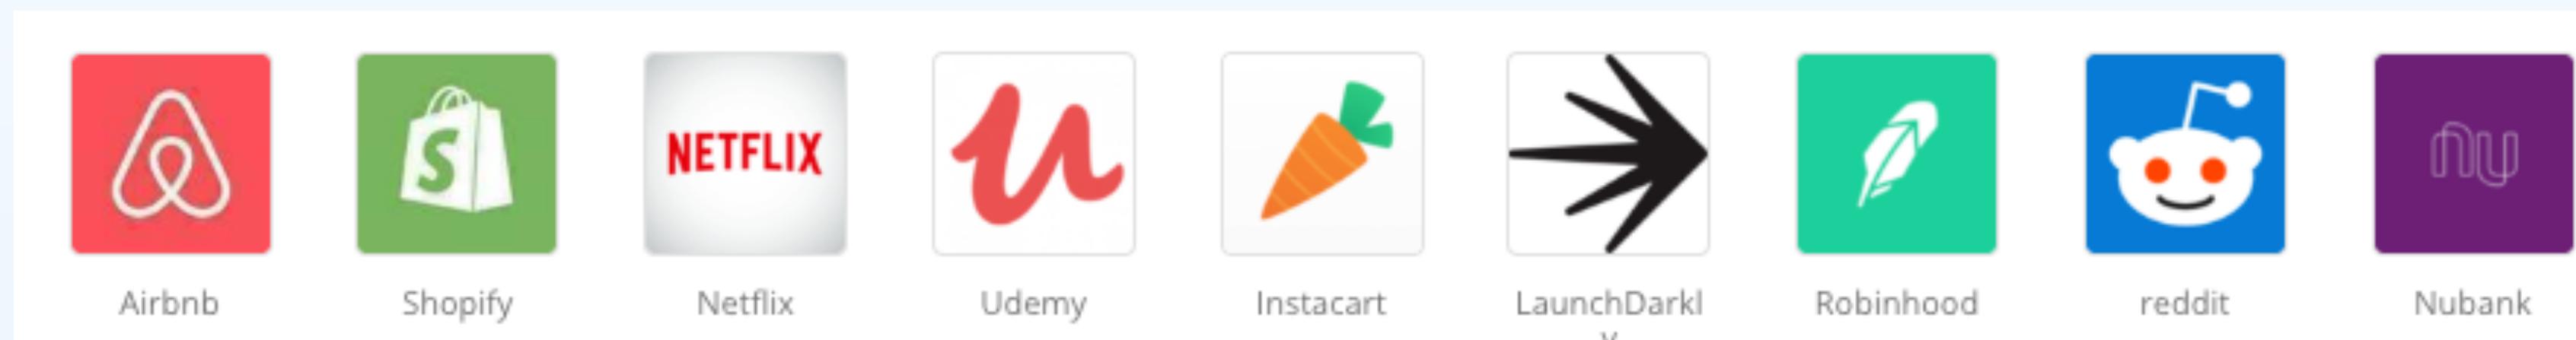
TypeScript, JavaScript, Python, C++, Go and Rust lead among top trending Github technology repos... while JavaScript, Python, and Java lead as primary languages among top trending Github educational repos



# Companies that use GitHub?

Stack share reports

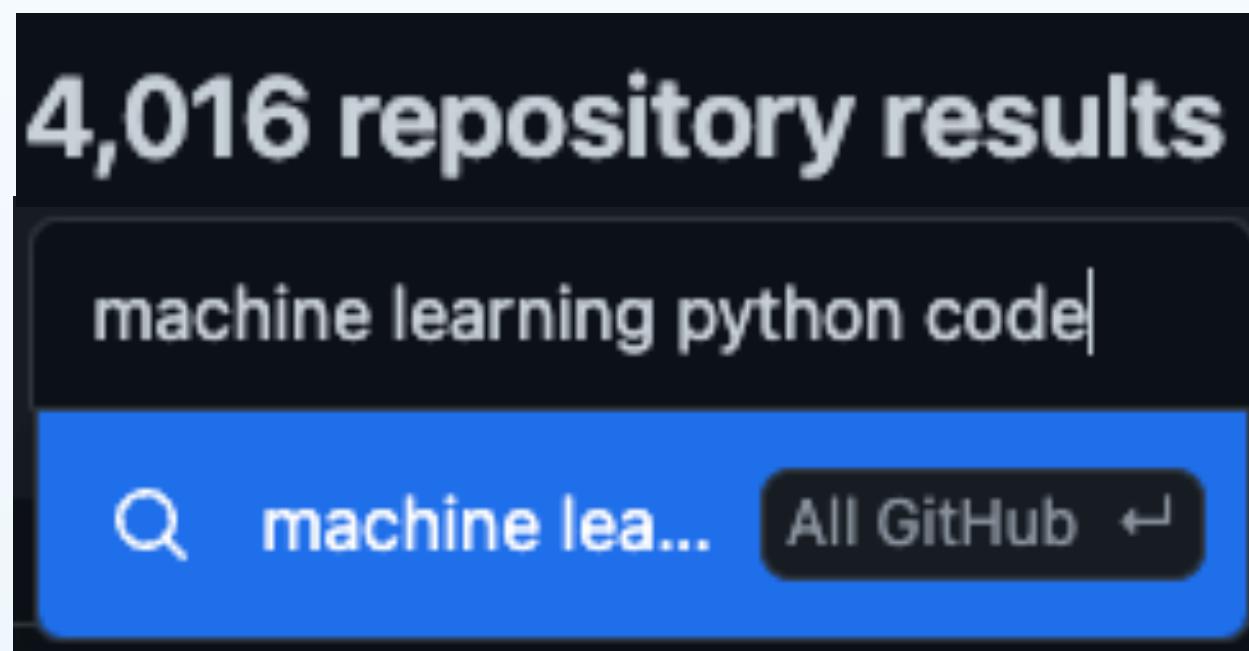
10029 companies reportedly use GitHub in their tech stacks, including Airbnb, Shopify, and Netflix.



# I.) Use other open source code

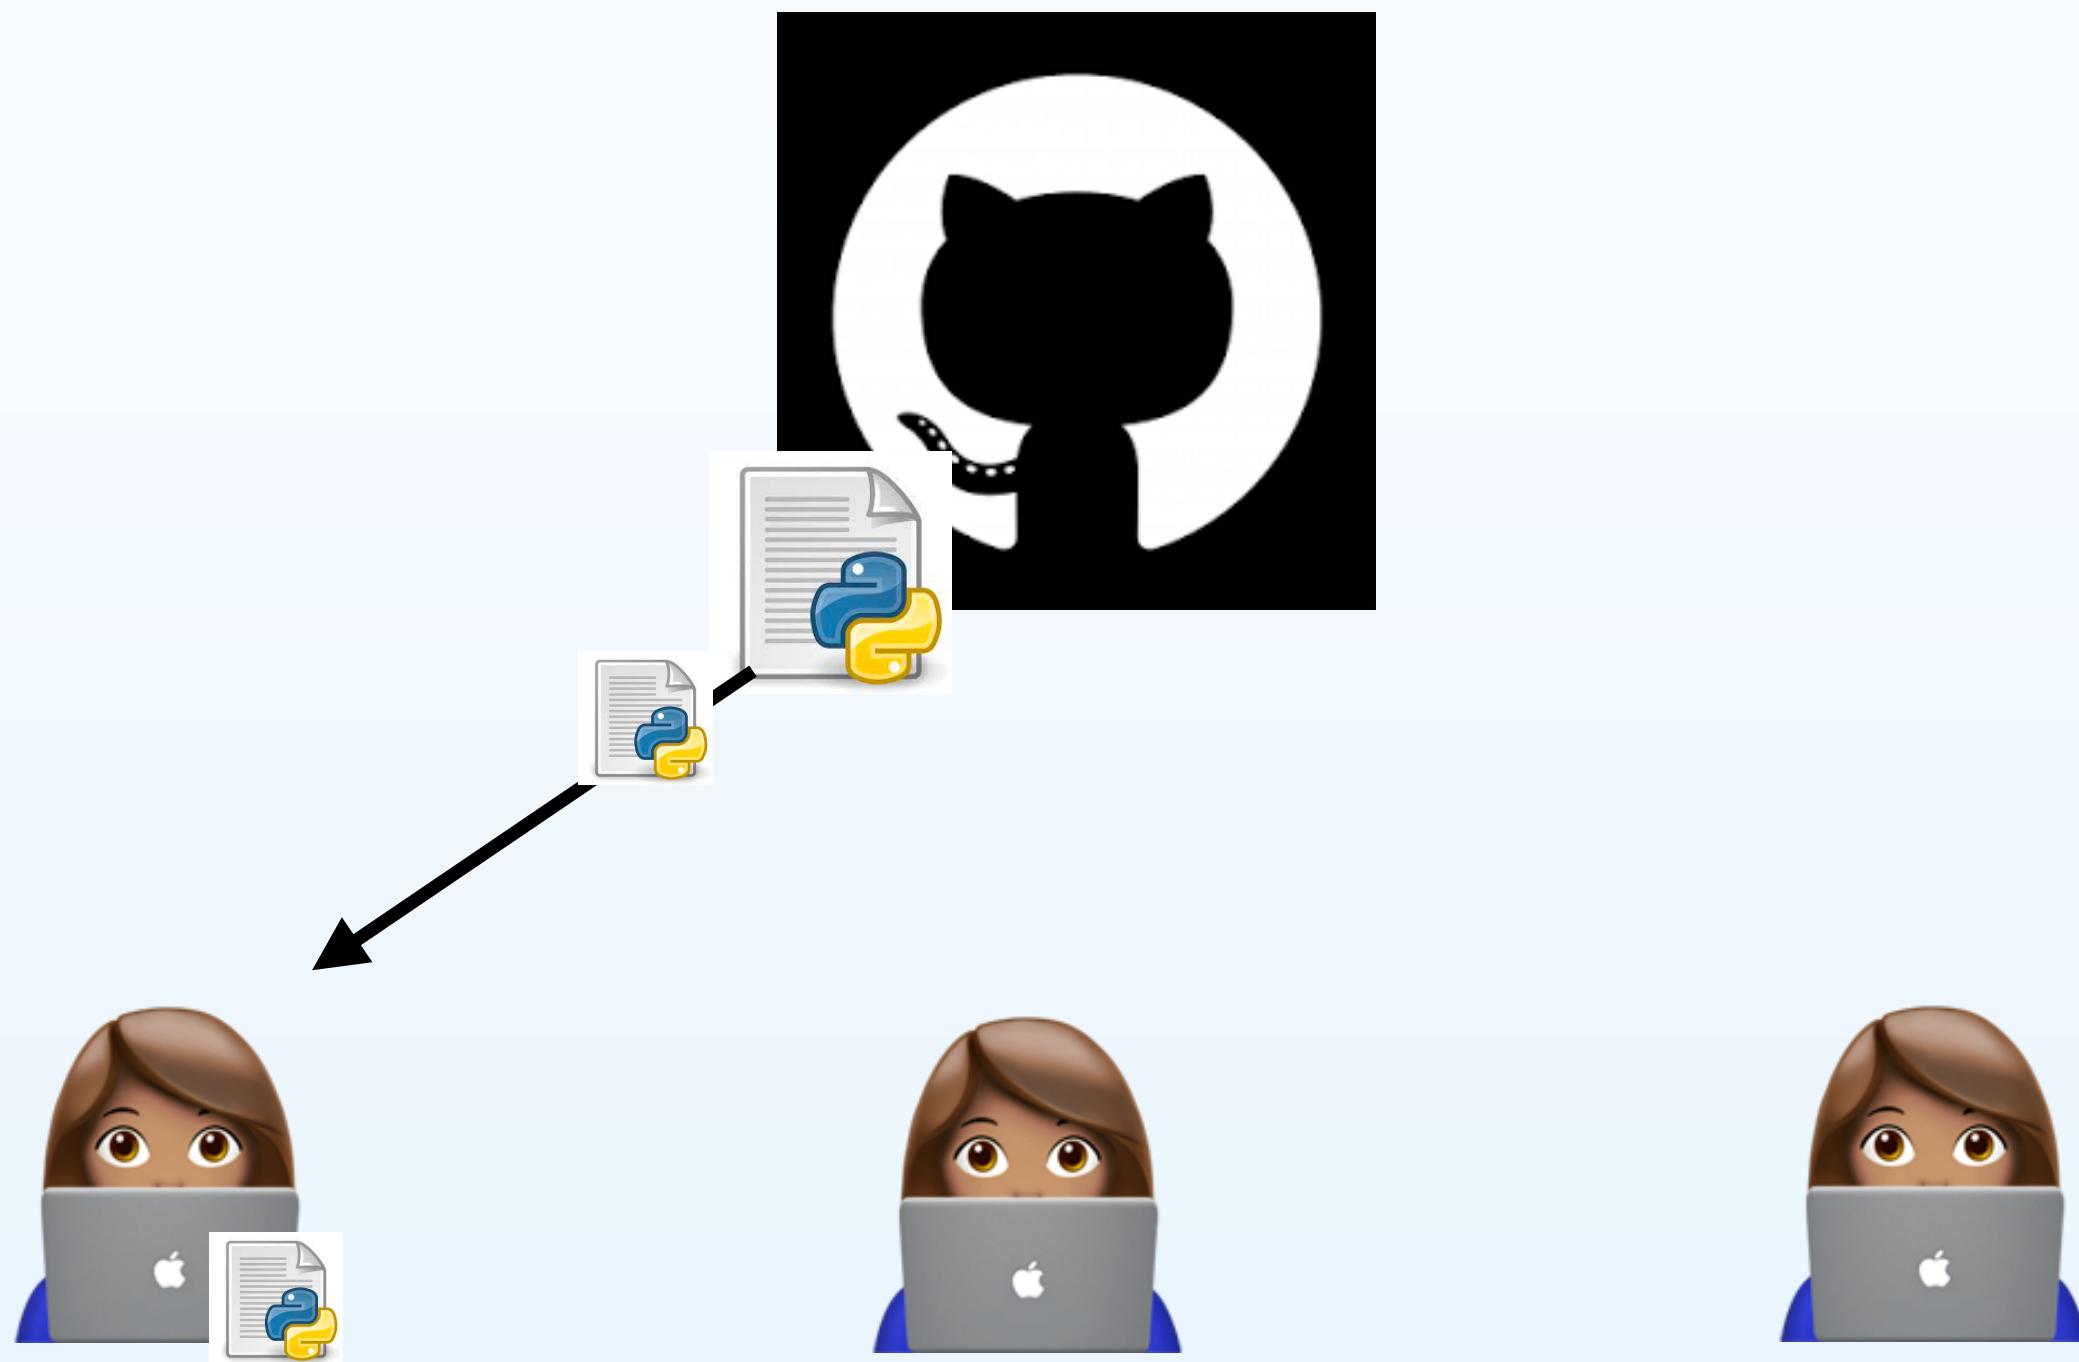
There is a huge amount of sample code you can learn from

This is an incredible resource



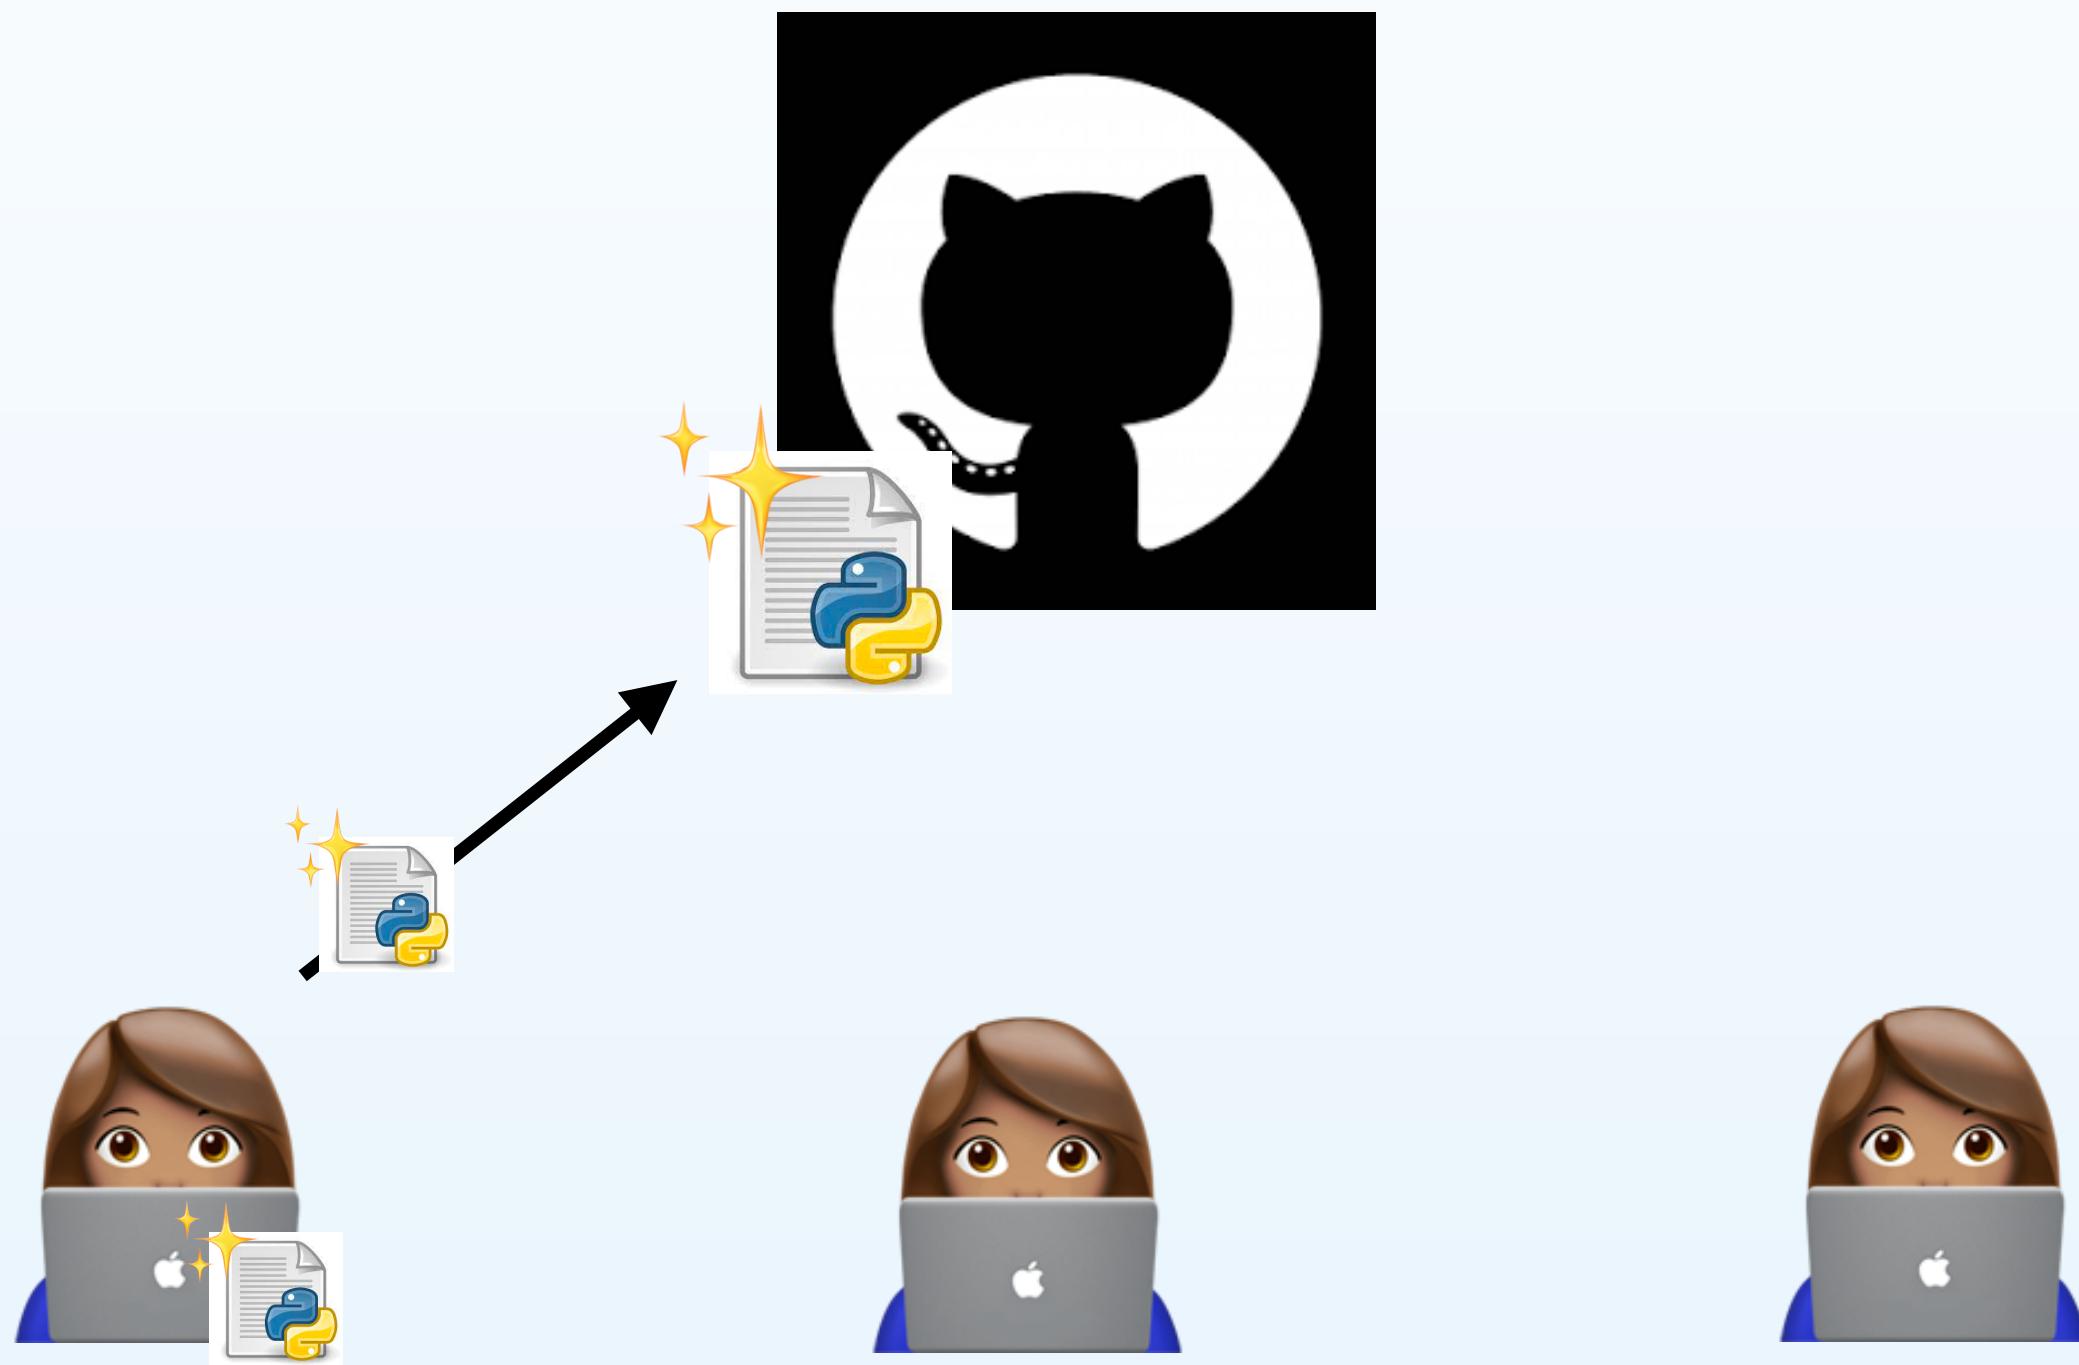
## 2.) Share code with a team

Github allows a team of data scientists to have a central location to pull the most recent code and post their most recent edits



## 2.) Share code with a team

Github allows a team of data scientists to have a central location to pull the most recent code and post their most recent edits

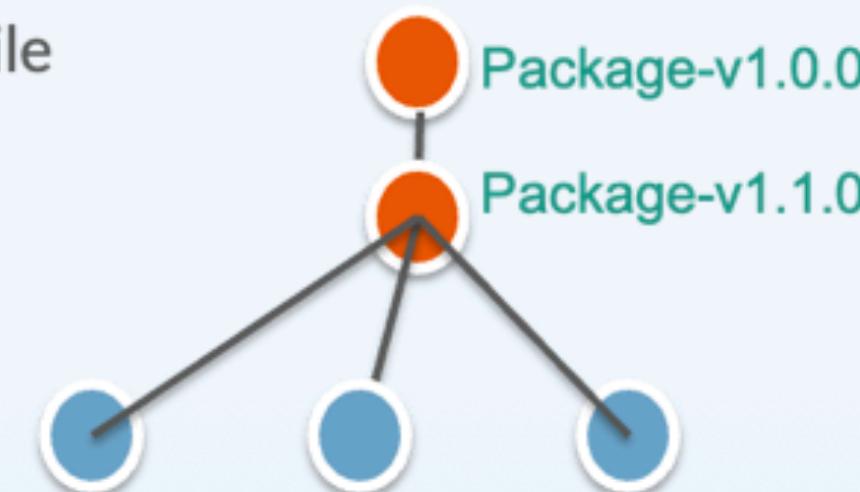


# 3.) Version Control

Version Control : Tracking changes in any set of files, usually used for coordinating work among programmers collaboratively

## Branching

Imagine 3 developers each working in an isolated copy of the working file



# 3.) Version Control

Version Control : Tracking changes in any set of files, usually used for coordinating work among programmers collaboratively

## Merge

Changes are then merged into the develop branch after quality check

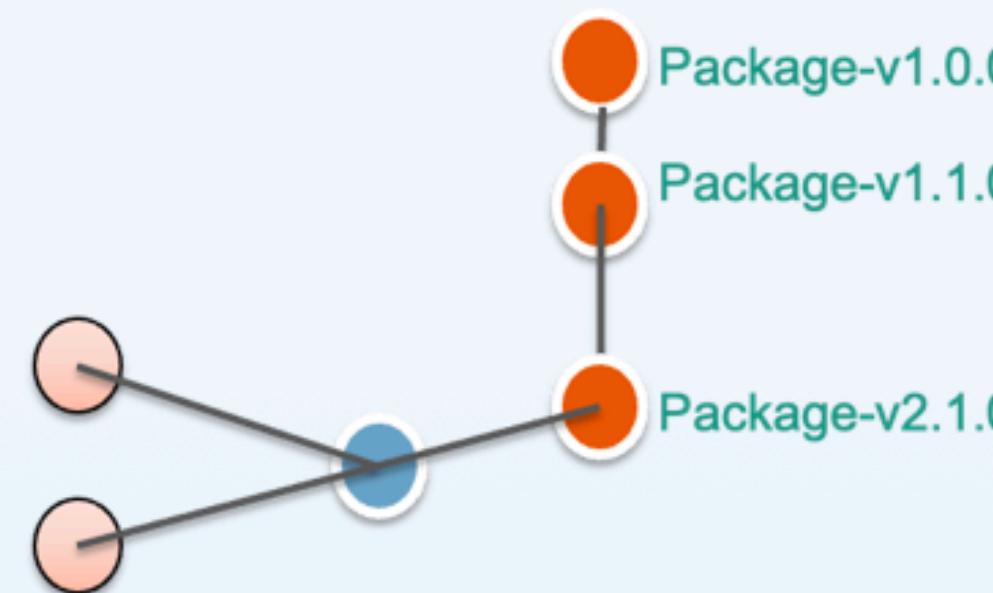


# 3.) Version Control

Version Control : Tracking changes in any set of files, usually used for coordinating work among programmers collaboratively

## Branch of a Branch

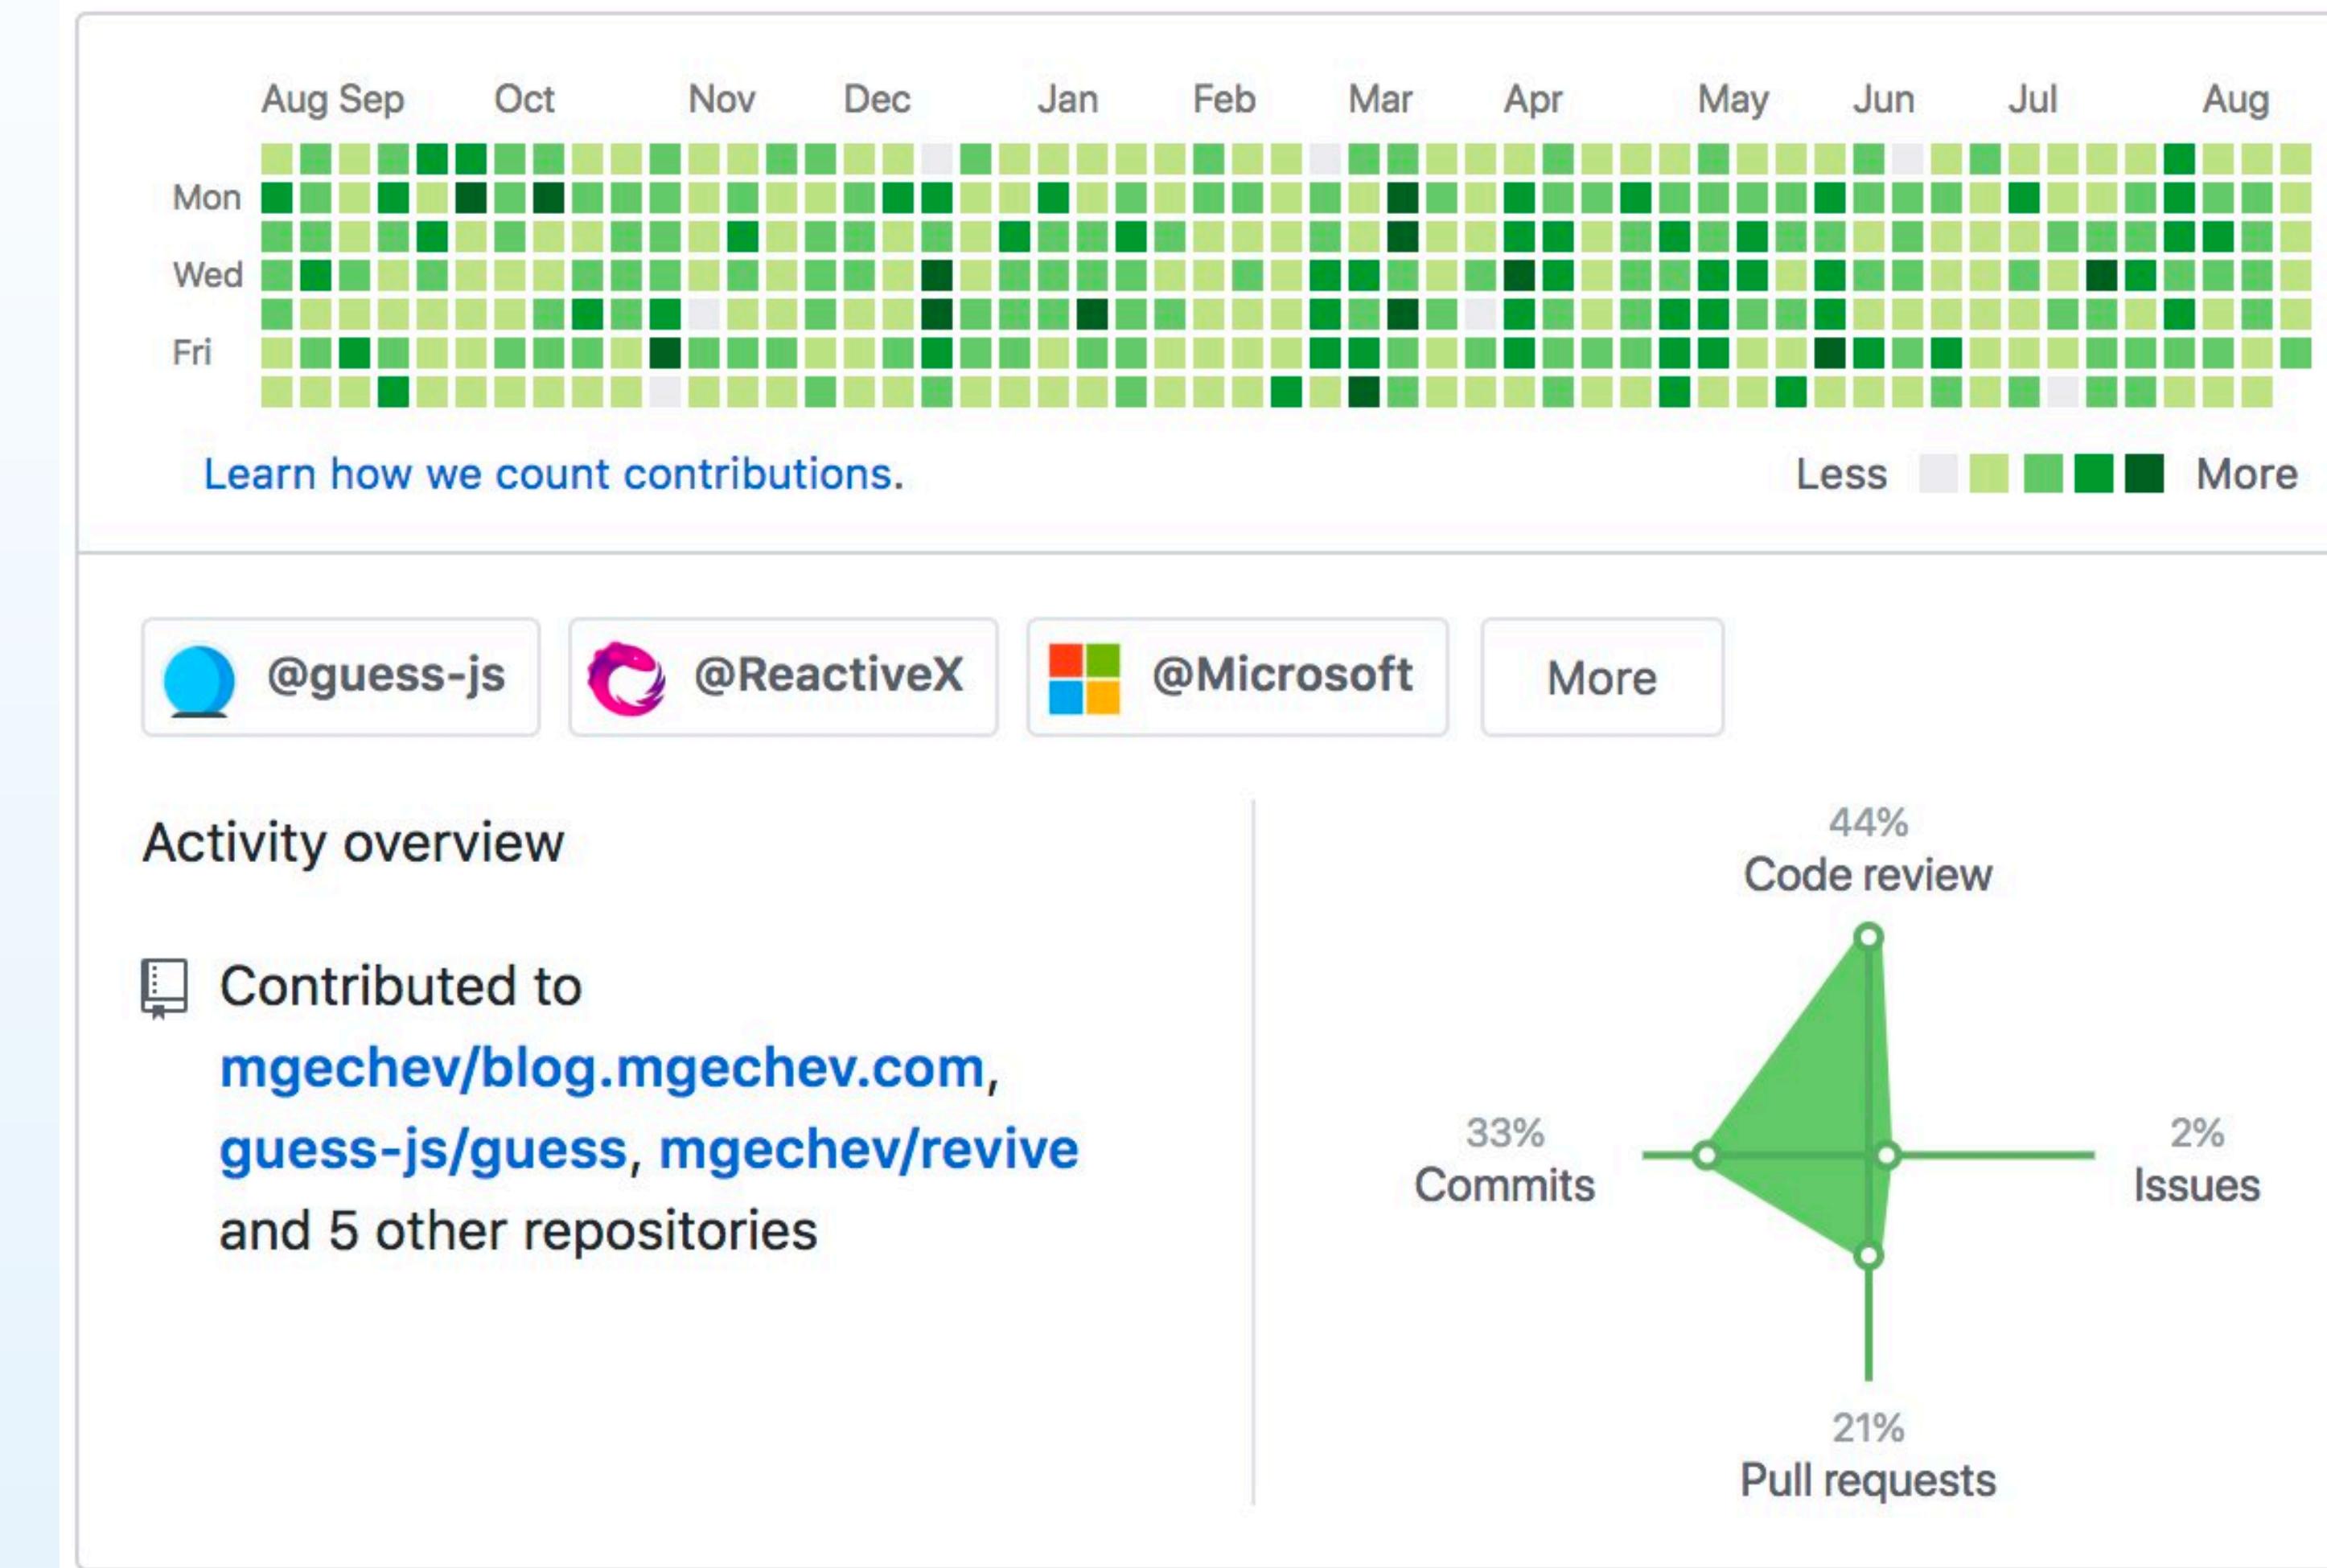
Can make branches of branches



# 4.) Signal to Companies you are an Active developer

More and more often for data driven roles, employers will request your GitHub

Store your projects here so they can see your work

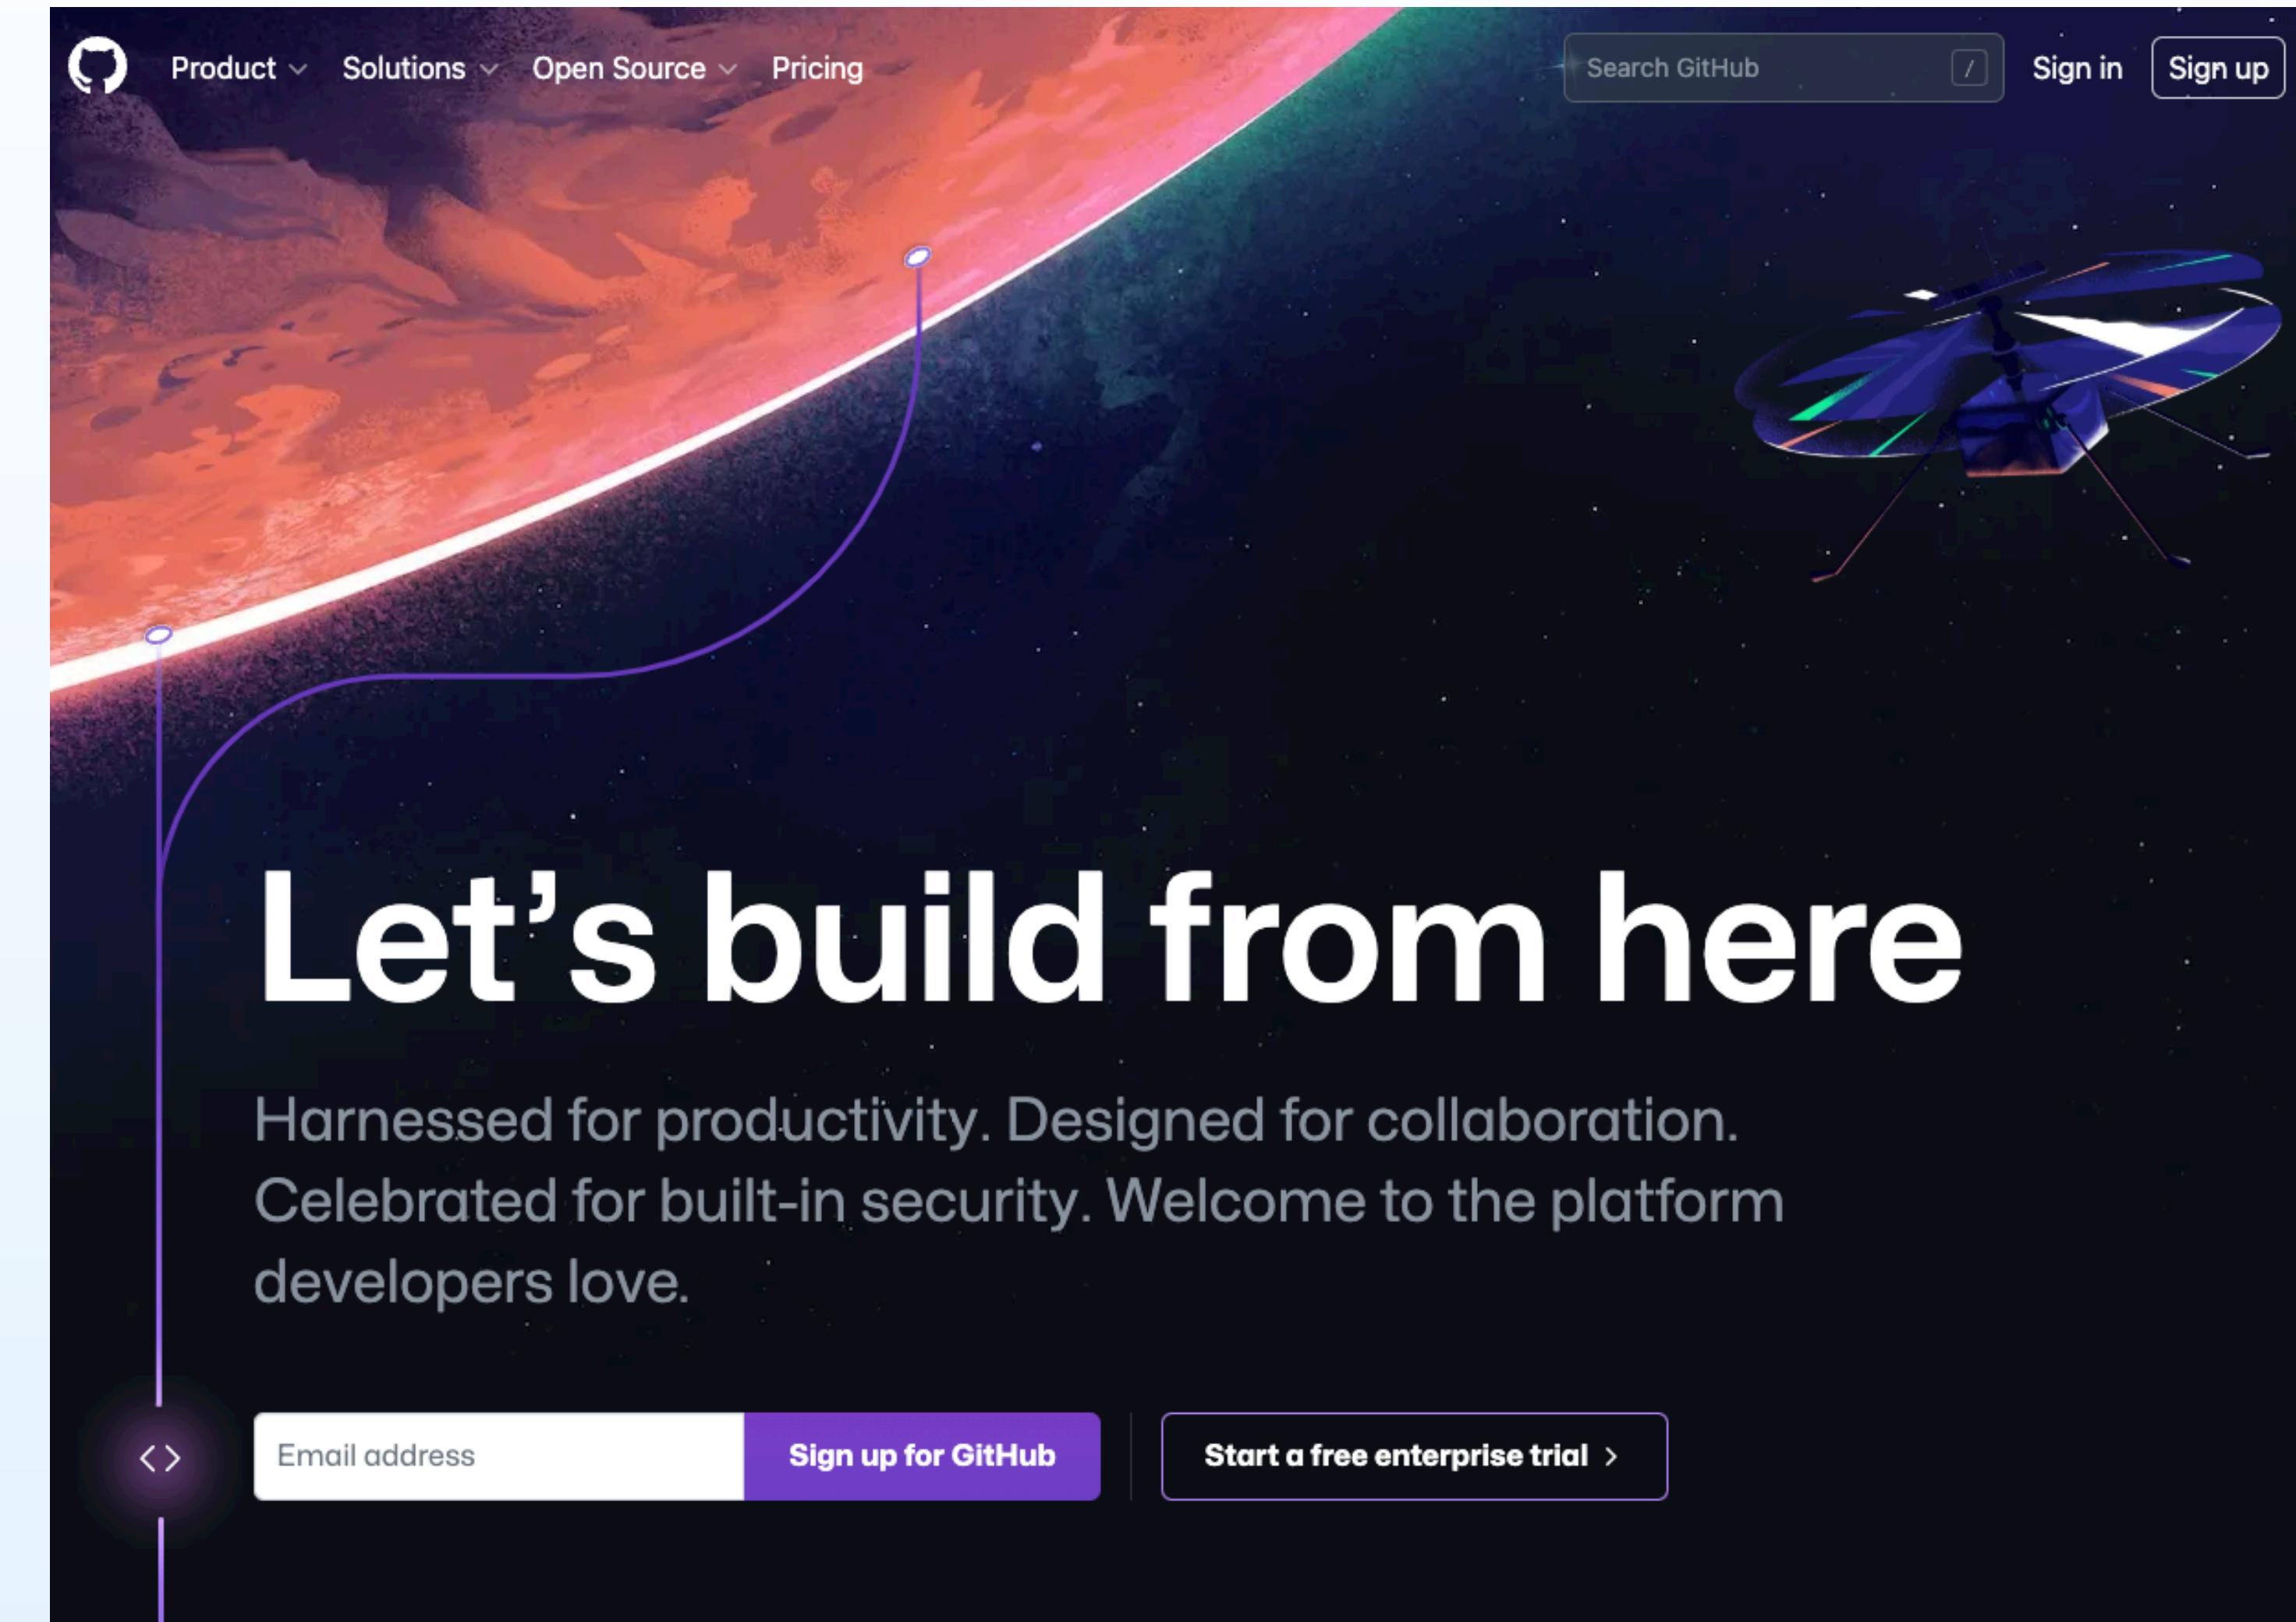


# Make a GitHub Account

[GitHub.com](https://github.com)

Username

InClassEcon441BDELETE



# Make a Public Repository

[GitHub.com](https://github.com)

**Start a new repository**

A repository contains all of your project's files, revision history, and collaborator discussion.

InClassEcon441BDELETE /

 **Public**  
Anyone on the internet can see this repository

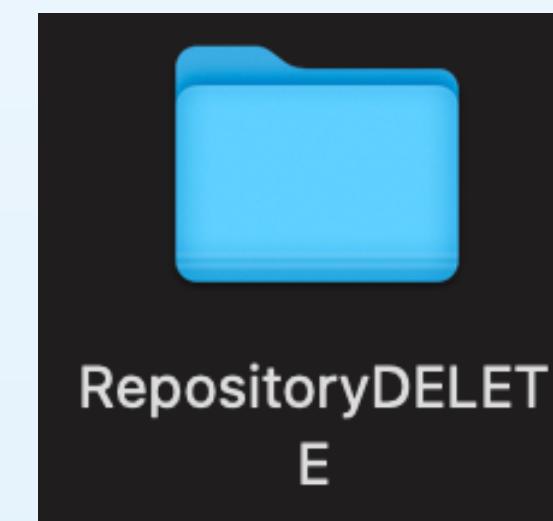
 **Private**  
You choose who can see and commit to this repository

**Create a new repository**

Name your repository

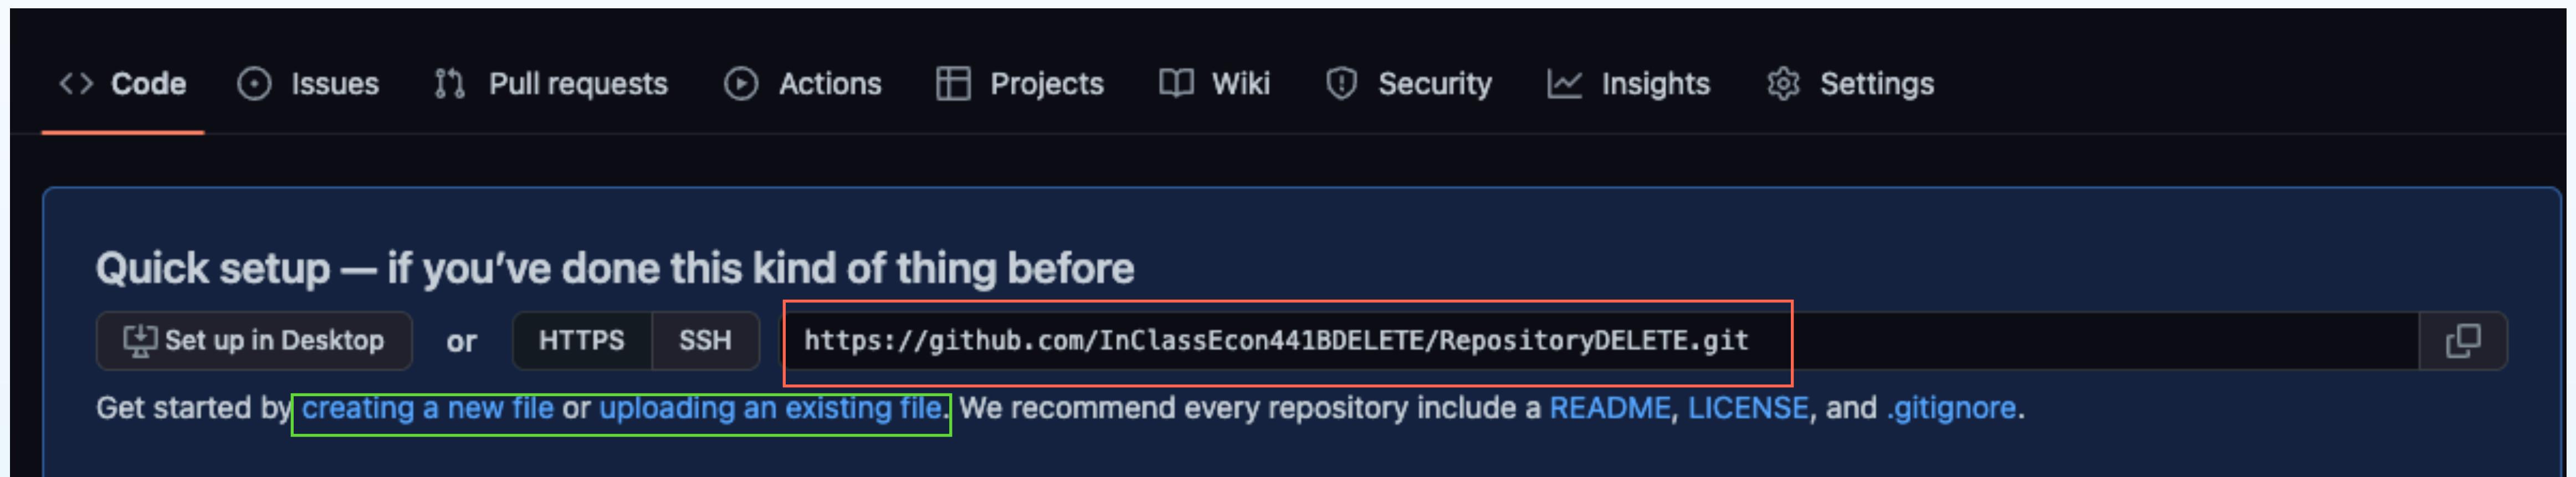
 **InClassEcon441BDELETE / RepositoryDELETE**  **Public**

Repositories work the same as a folder on your computer



# Adding to the repository

There are several ways to get code into a repository

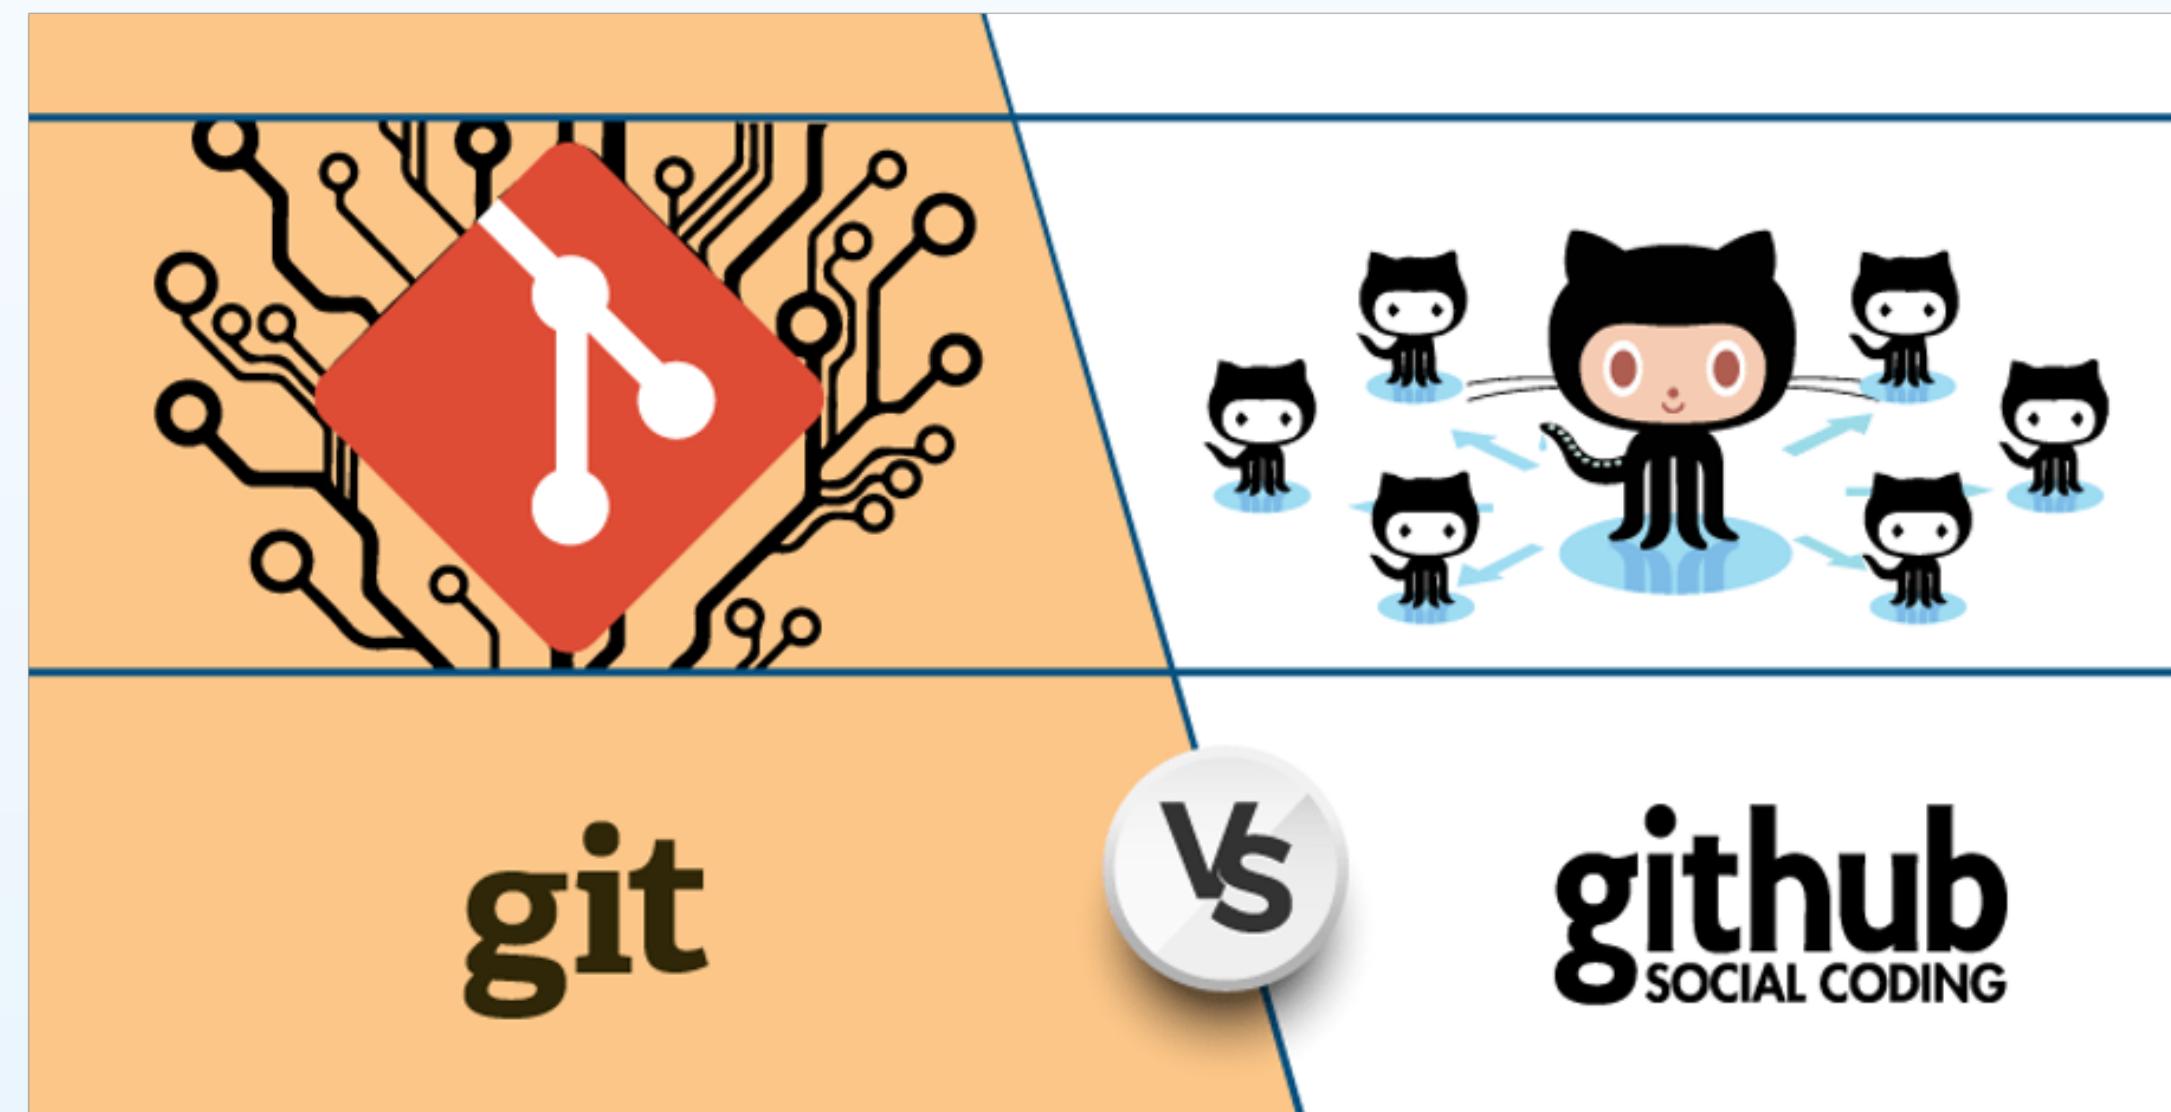


This is what you will submit to Bruin Learn at the end of class

This is how we will begin to upload files

# What is Git?

Is the underlying code that Github lets visualize



Git is a version control system  
Keeps track of your changes to code

A cloud based interface that  
Helps visualize what Git is doing

# What is Git?

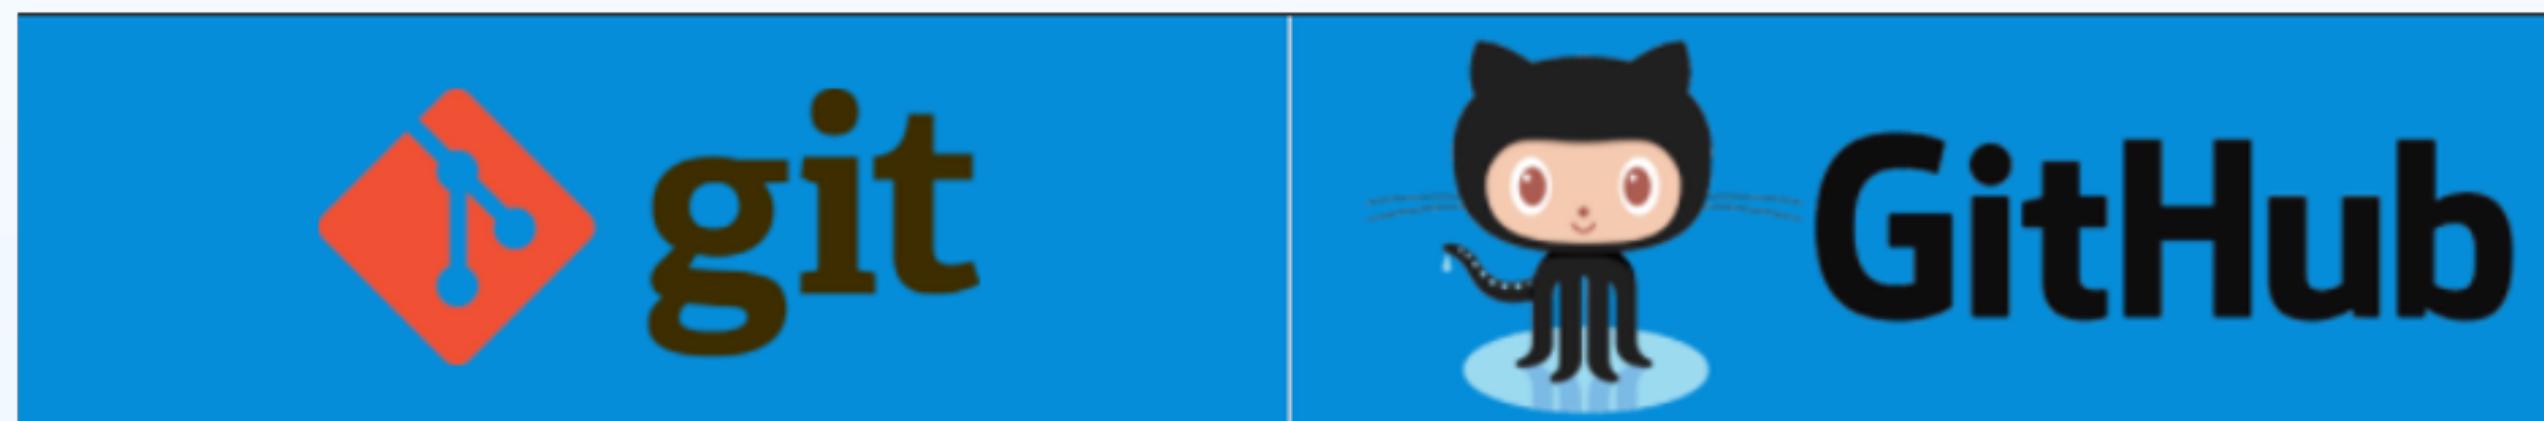
It is a software, on your local computer that allows you to do the version controlling and posting to Github  
You can use GitHub with or without using Git

|  git |  GitHub                     |
|---|--|
| 1. It is a software   | 1. It is a service   |
| 2. It is installed locally on the system  | 2. It is hosted on Web   |
| 3. It is a command line tool  | 3. It provides a graphical interface   |
| 4. It is a tool to manage different versions of edits, made to files in a git repository  | 4. It is a space to upload a copy of the <b>Git</b> repository   |
| 5. It provides functionalities like Version Control System Source Code Management         | 5. It provides functionalities of Git like VCS, Source Code Management as well as adding few of its own features |

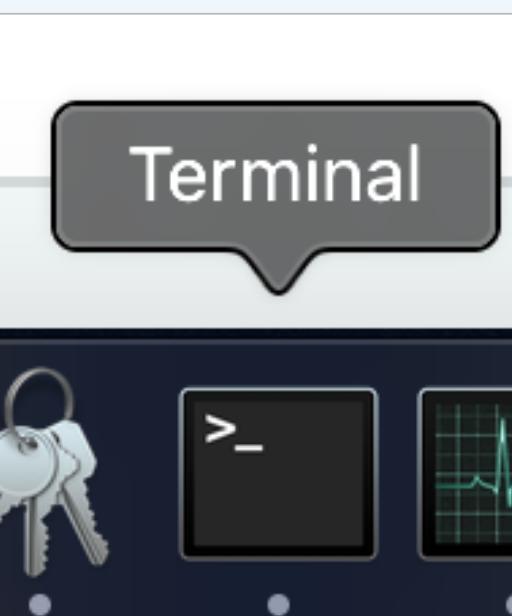
# What is Git?

Git uses commands written in shell scripts

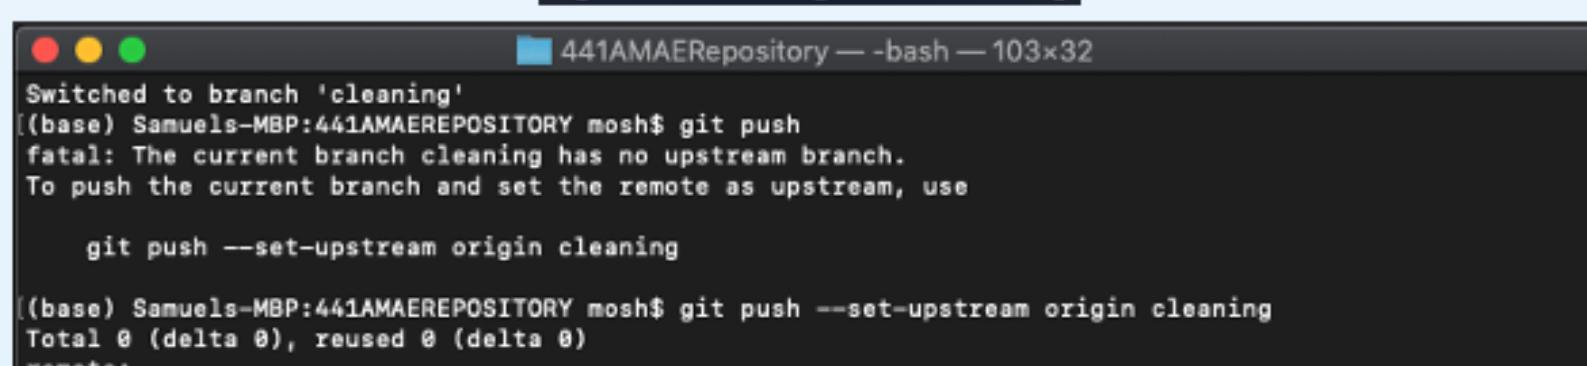
Accessible through



Terminal



A screenshot of a Mac OS X Dock. A speech bubble icon contains the word "Terminal". Below it, there are icons for a keychain, a terminal window with a prompt, and a graph.



A screenshot of a macOS Terminal window. The title bar says "441AMAERepository — bash — 103x32". The window displays the following text:

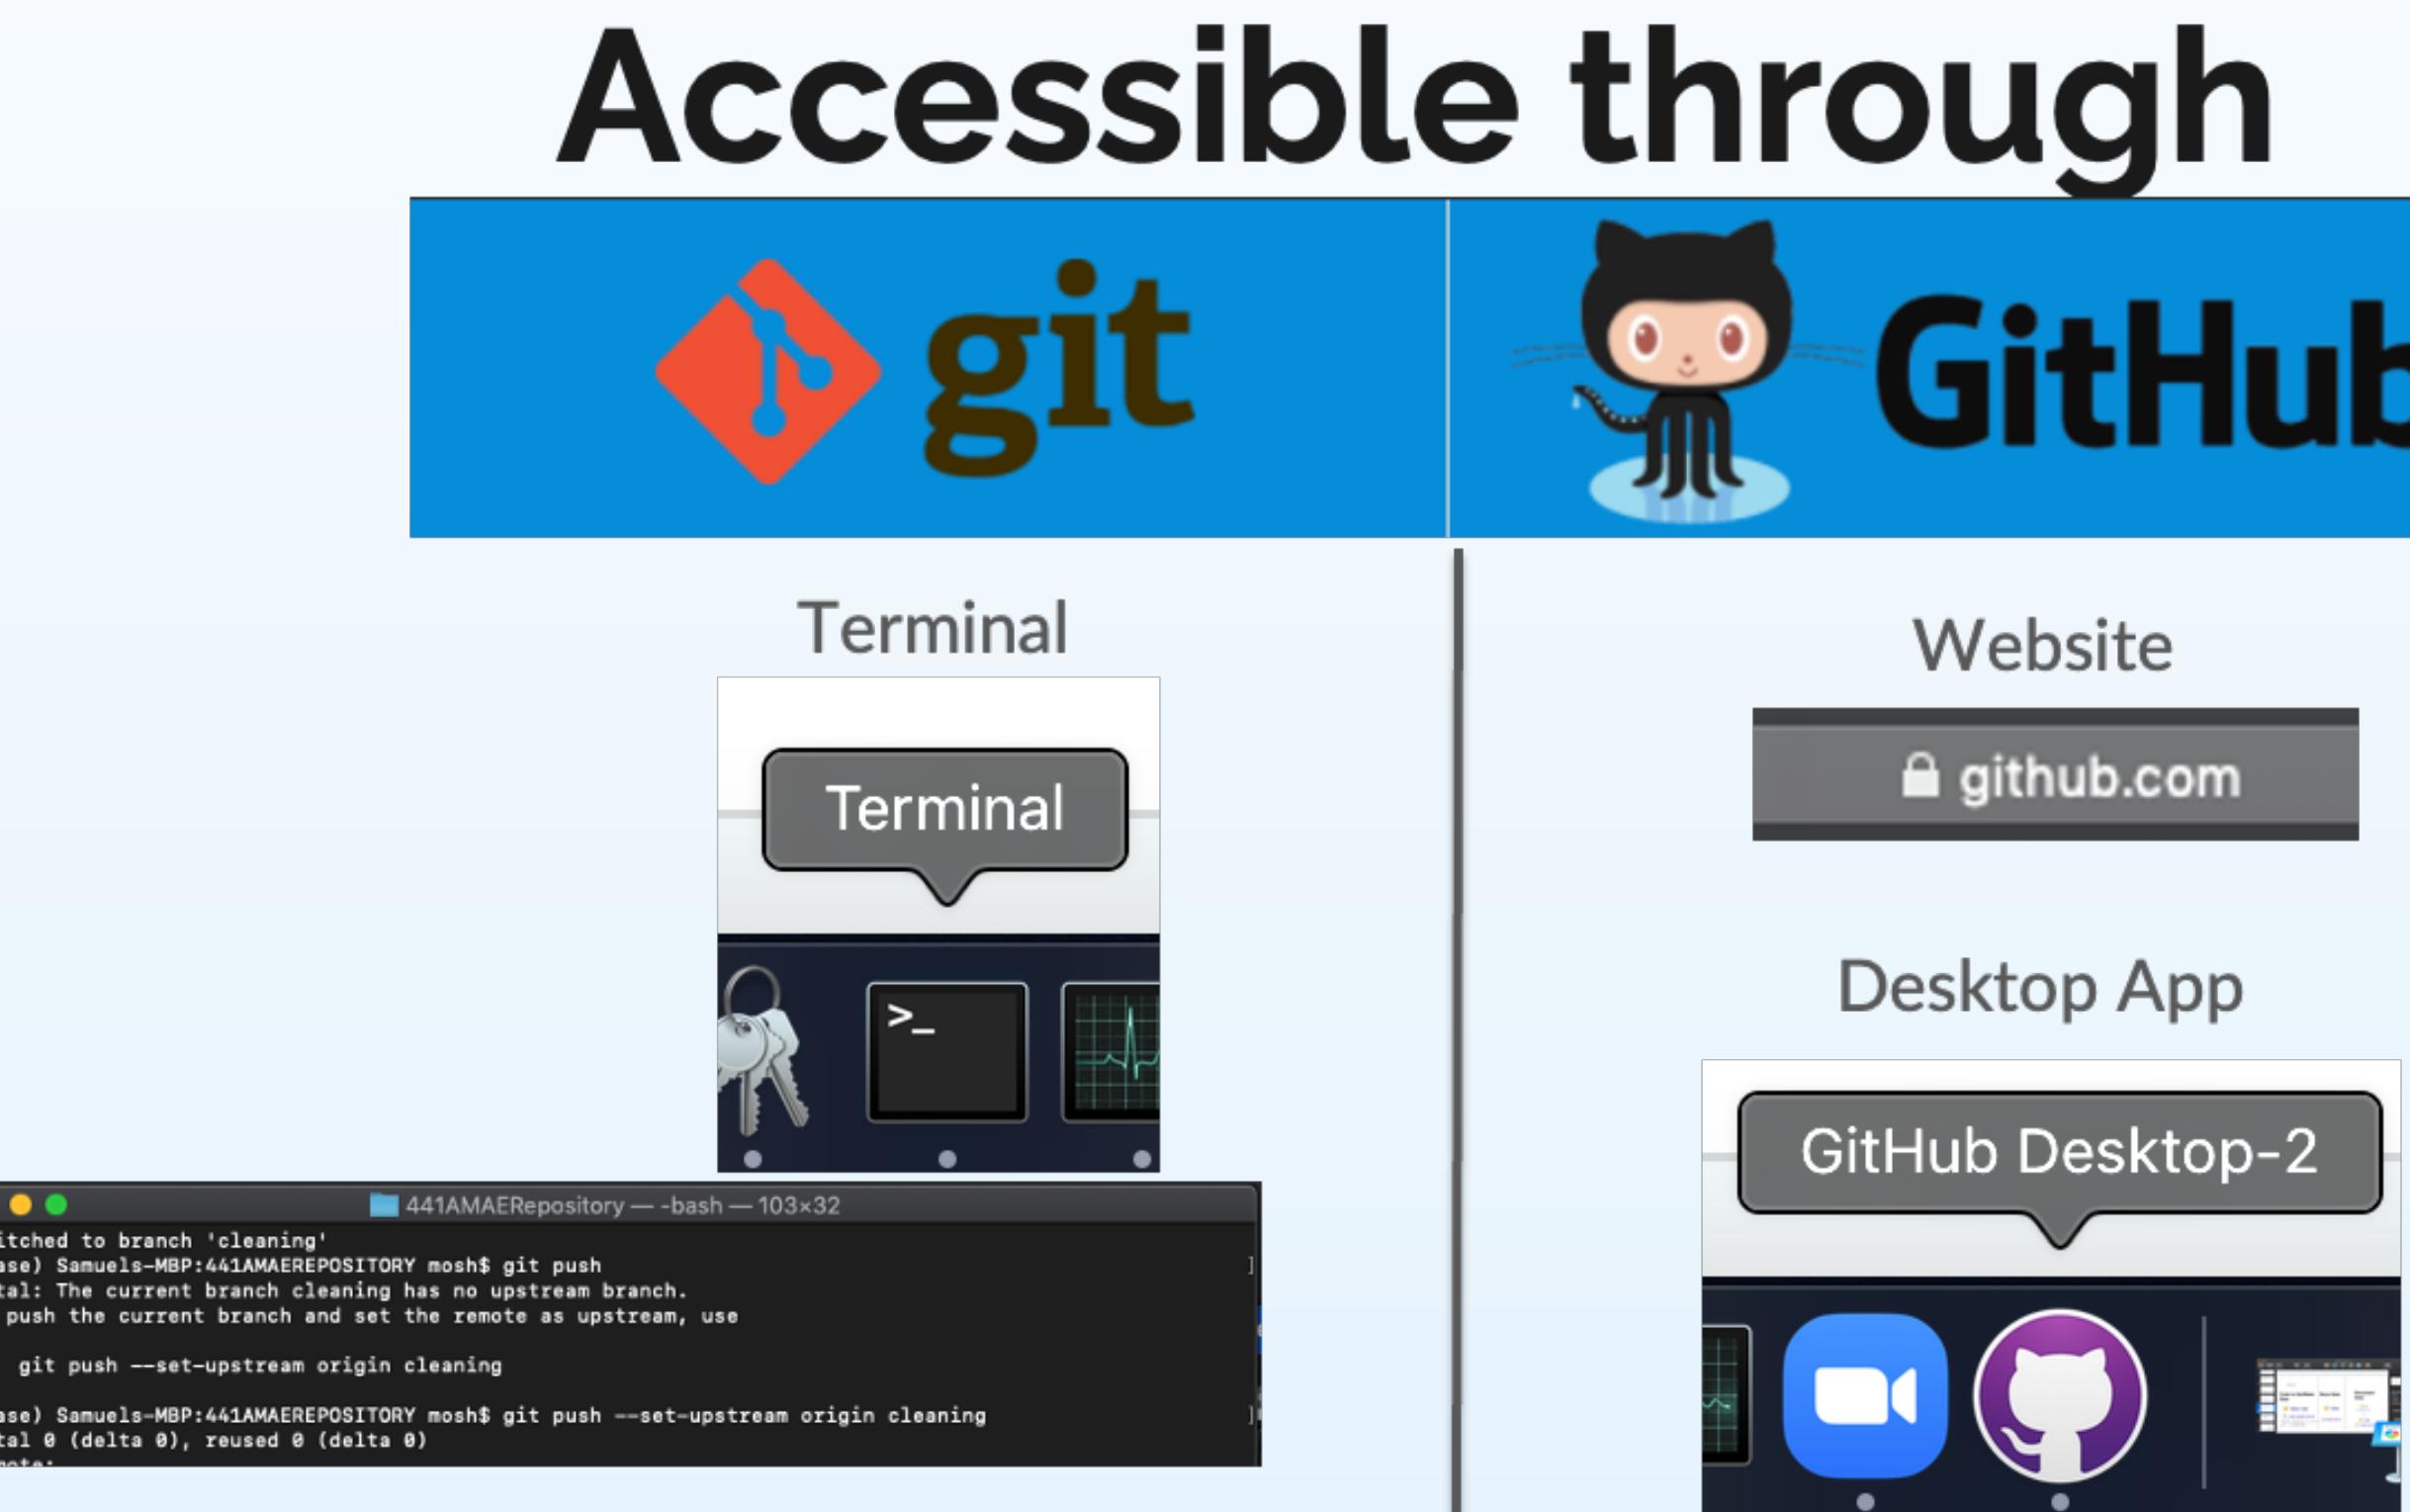
```
Switched to branch 'cleaning'
(base) Samuels-MBP:441AMAEREPOSITORY mosh$ git push
fatal: The current branch cleaning has no upstream branch.
To push the current branch and set the remote as upstream, use

    git push --set-upstream origin cleaning

(base) Samuels-MBP:441AMAEREPOSITORY mosh$ git push --set-upstream origin cleaning
Total 0 (delta 0), reused 0 (delta 0)
remote:
```

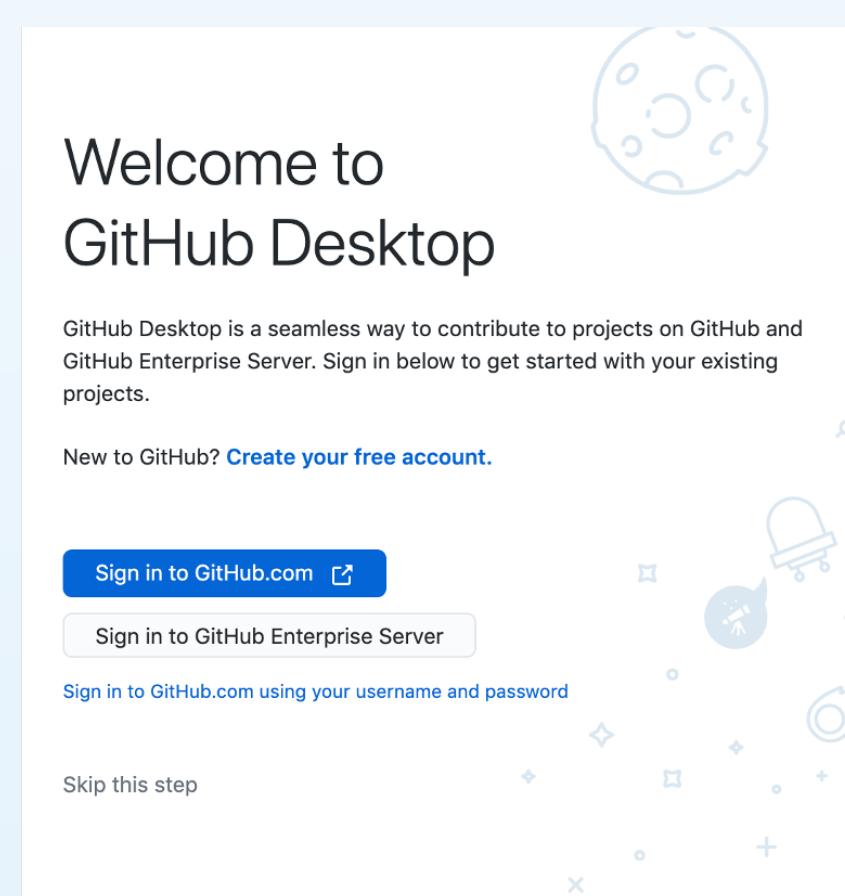
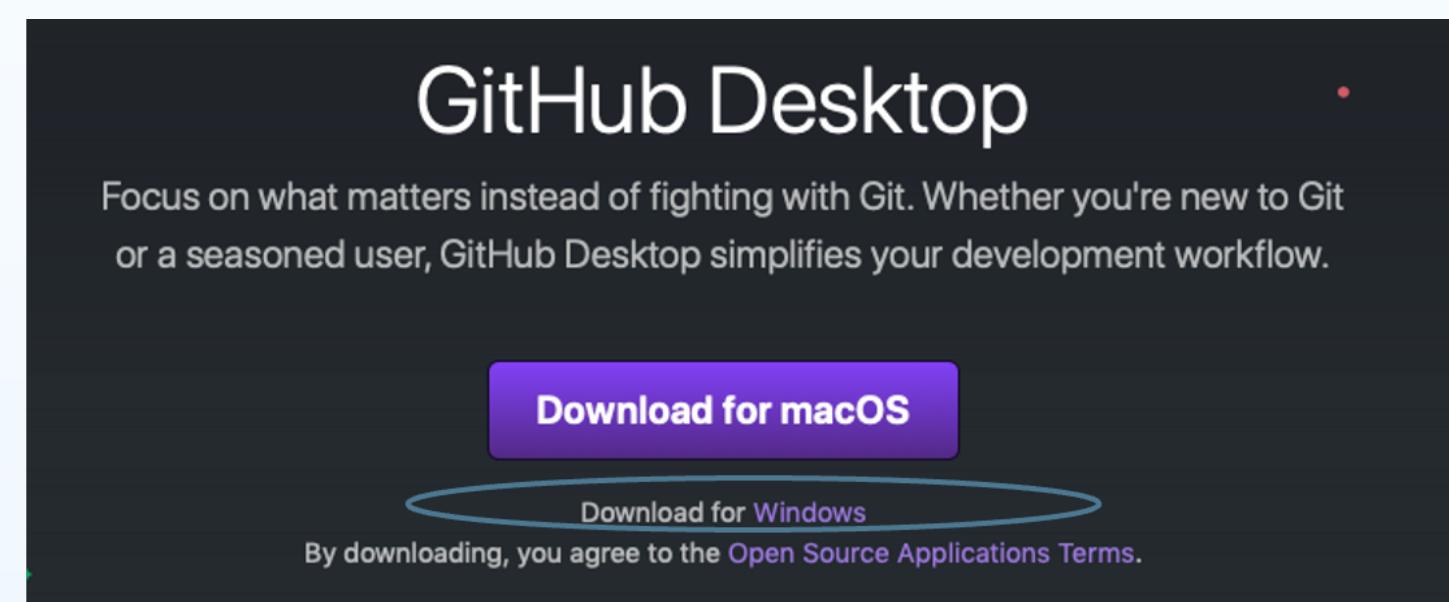
# What is Git?

Github can be used without writing any code



# There is a lot we do not cover about Git/Github

If this is something that interests you: Research and practice more

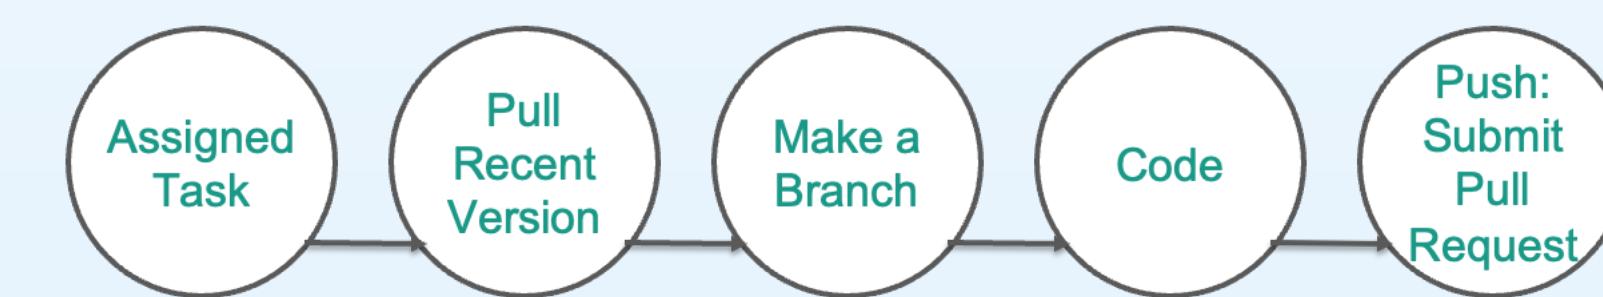


- Step 1: Download Github desktop**
- Step 2: Make a Github account**
- Step 3: Make a Repository**
  - 3.a.) Posting on the Website
  - 3.b.) Clone a repository
  - 3.c.) Create a new repository
  - 3.d.) Add existing Repository
- Step 4: Get coding/editing**

**Team Member validates work and merges**

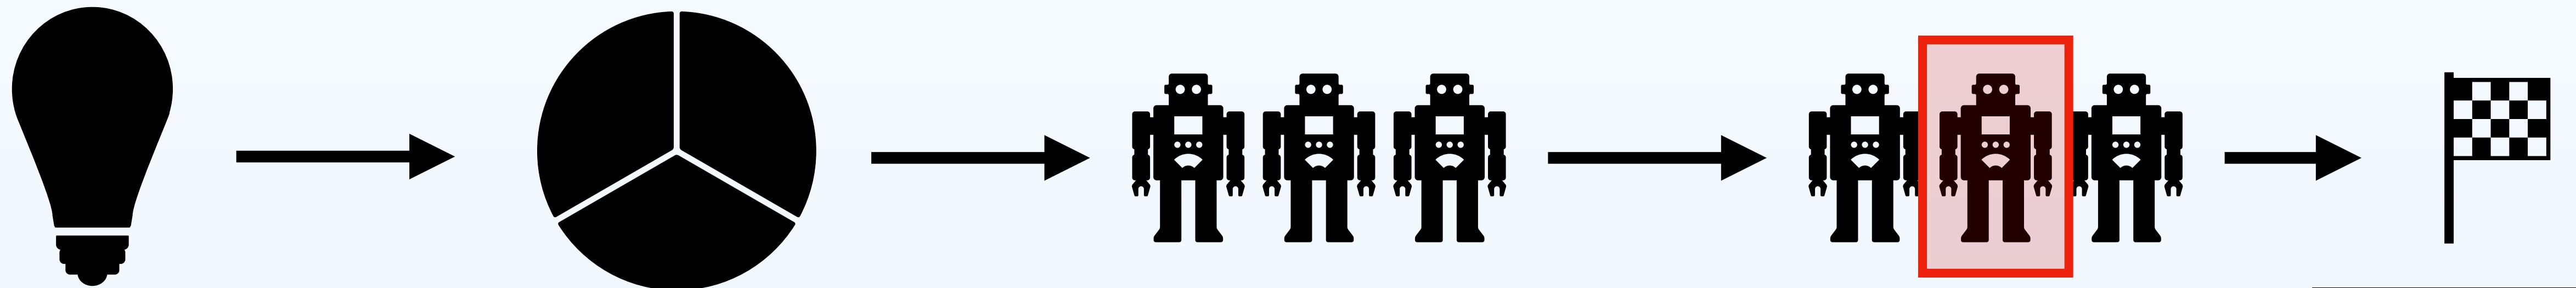
A screenshot of a GitHub pull request merge interface. It shows a green "Merge pull request" button. To the right, a note says: "You can also open this in GitHub Desktop or view command line instructions." Below this, a message says: "Pull request successfully merged and closed" and "You're all set—the lasttest branch can be safely deleted." There is also a "Delete branch" button.

## Work Flow



# Robust Model Building

# Machine Learning Pipeline



Have an Assumption :  
This data has predictive  
capabilities over the data

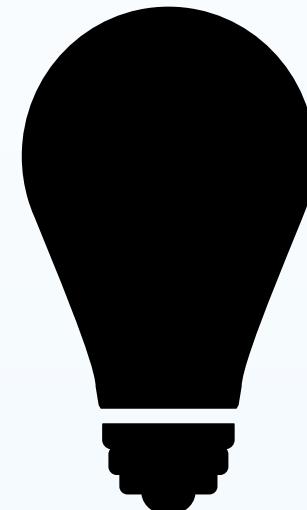
Data is randomly split into  
3 parts :  
In-sample, Out of Sample  
and Hold out

Various models are built  
using In-Sample Data with  
data best practices such  
as Cross Validation

Final Model is selected  
based on Out-Of Sample  
performance

If the final model  
does not perform  
as expected on  
Hold-Out  
Datapractices  
such as Cross  
Validation. The  
assumption is  
thrown out

# Machine Learning Pipeline : Example Regression



Have an Assumption :  
This data has predictive  
capabilities over the data

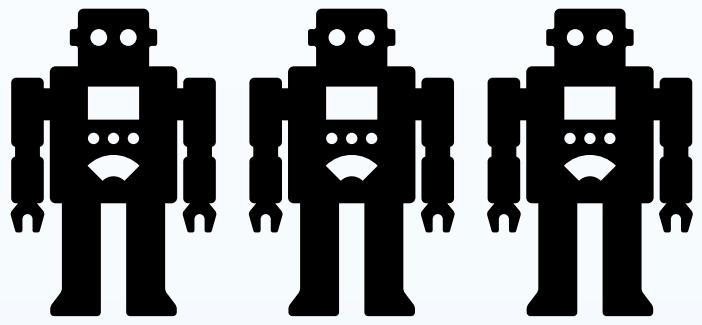


Data is randomly split into  
3 parts :  
In-sample, Out of Sample  
and Hold out

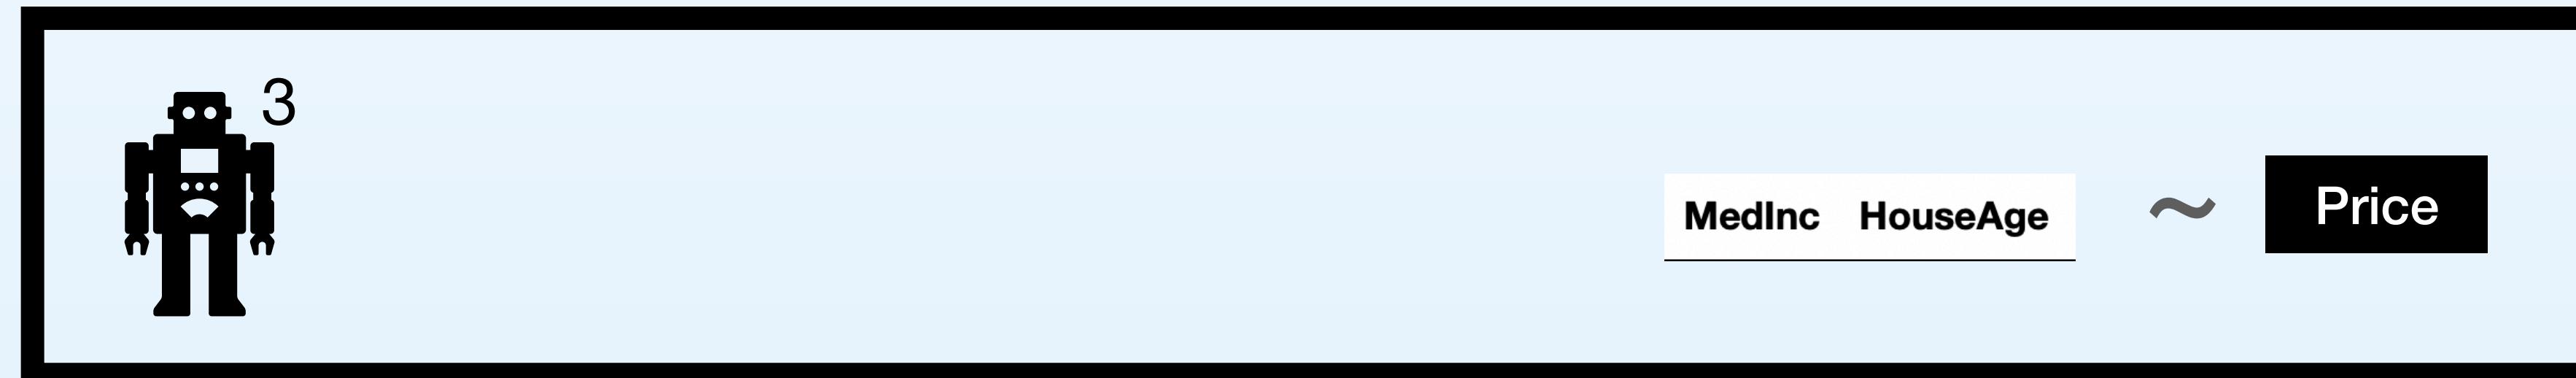
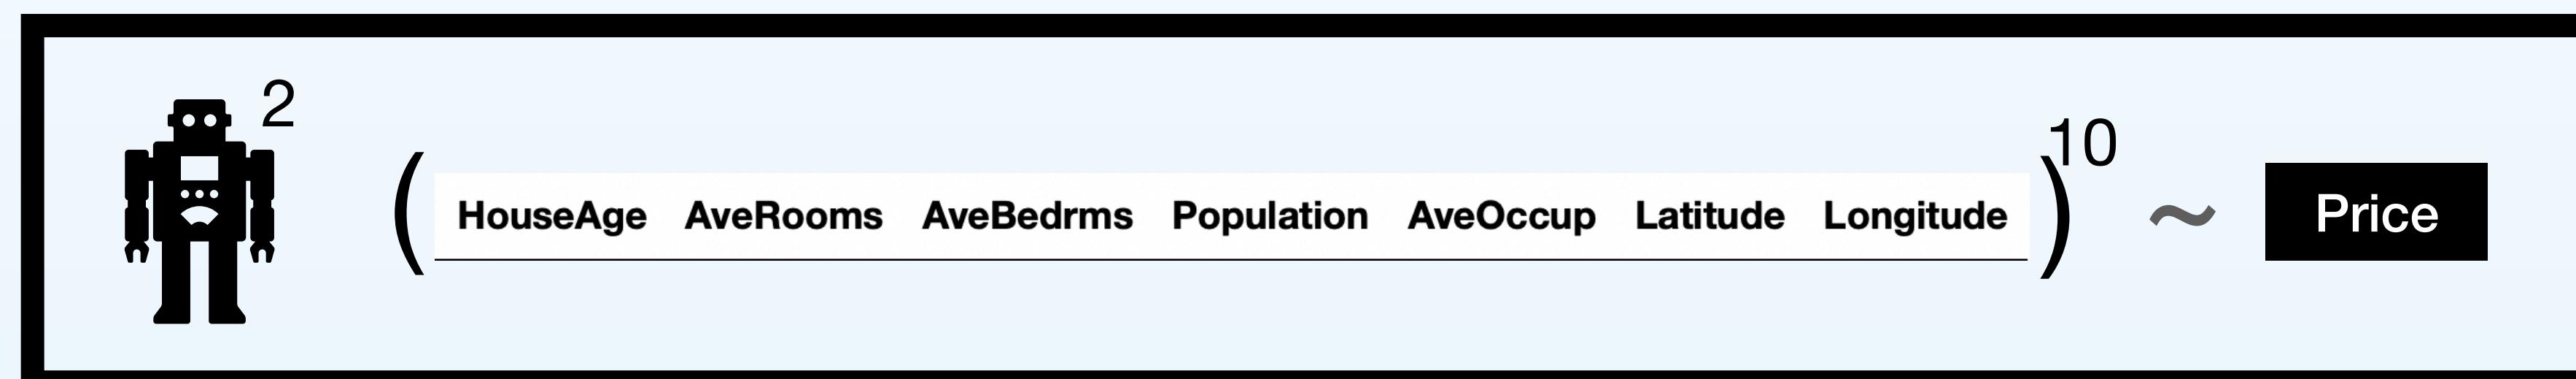
| MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | Longitude |
|--------|----------|----------|-----------|------------|----------|----------|-----------|
| 8.3252 | 41.0     | 6.984127 | 1.023810  | 322.0      | 2.555556 | 37.88    | -122.23   |
| 8.3014 | 21.0     | 6.238137 | 0.971880  | 2401.0     | 2.109842 | 37.86    | -122.22   |
| 7.2574 | 52.0     | 8.288136 | 1.073446  | 496.0      | 2.802260 | 37.85    | -122.24   |
| 5.6431 | 52.0     | 5.817352 | 1.073059  | 558.0      | 2.547945 | 37.85    | -122.25   |
| 3.8462 | 52.0     | 6.281853 | 1.081081  | 565.0      | 2.181467 | 37.85    | -122.25   |

This Data Predicts Price

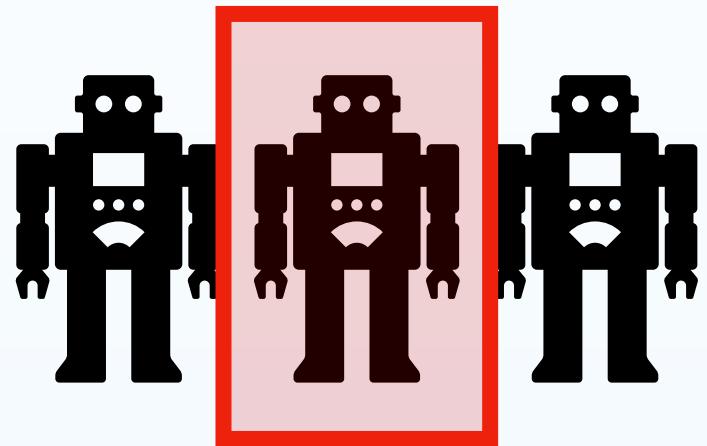
# Machine Learning Pipeline : Example Regression



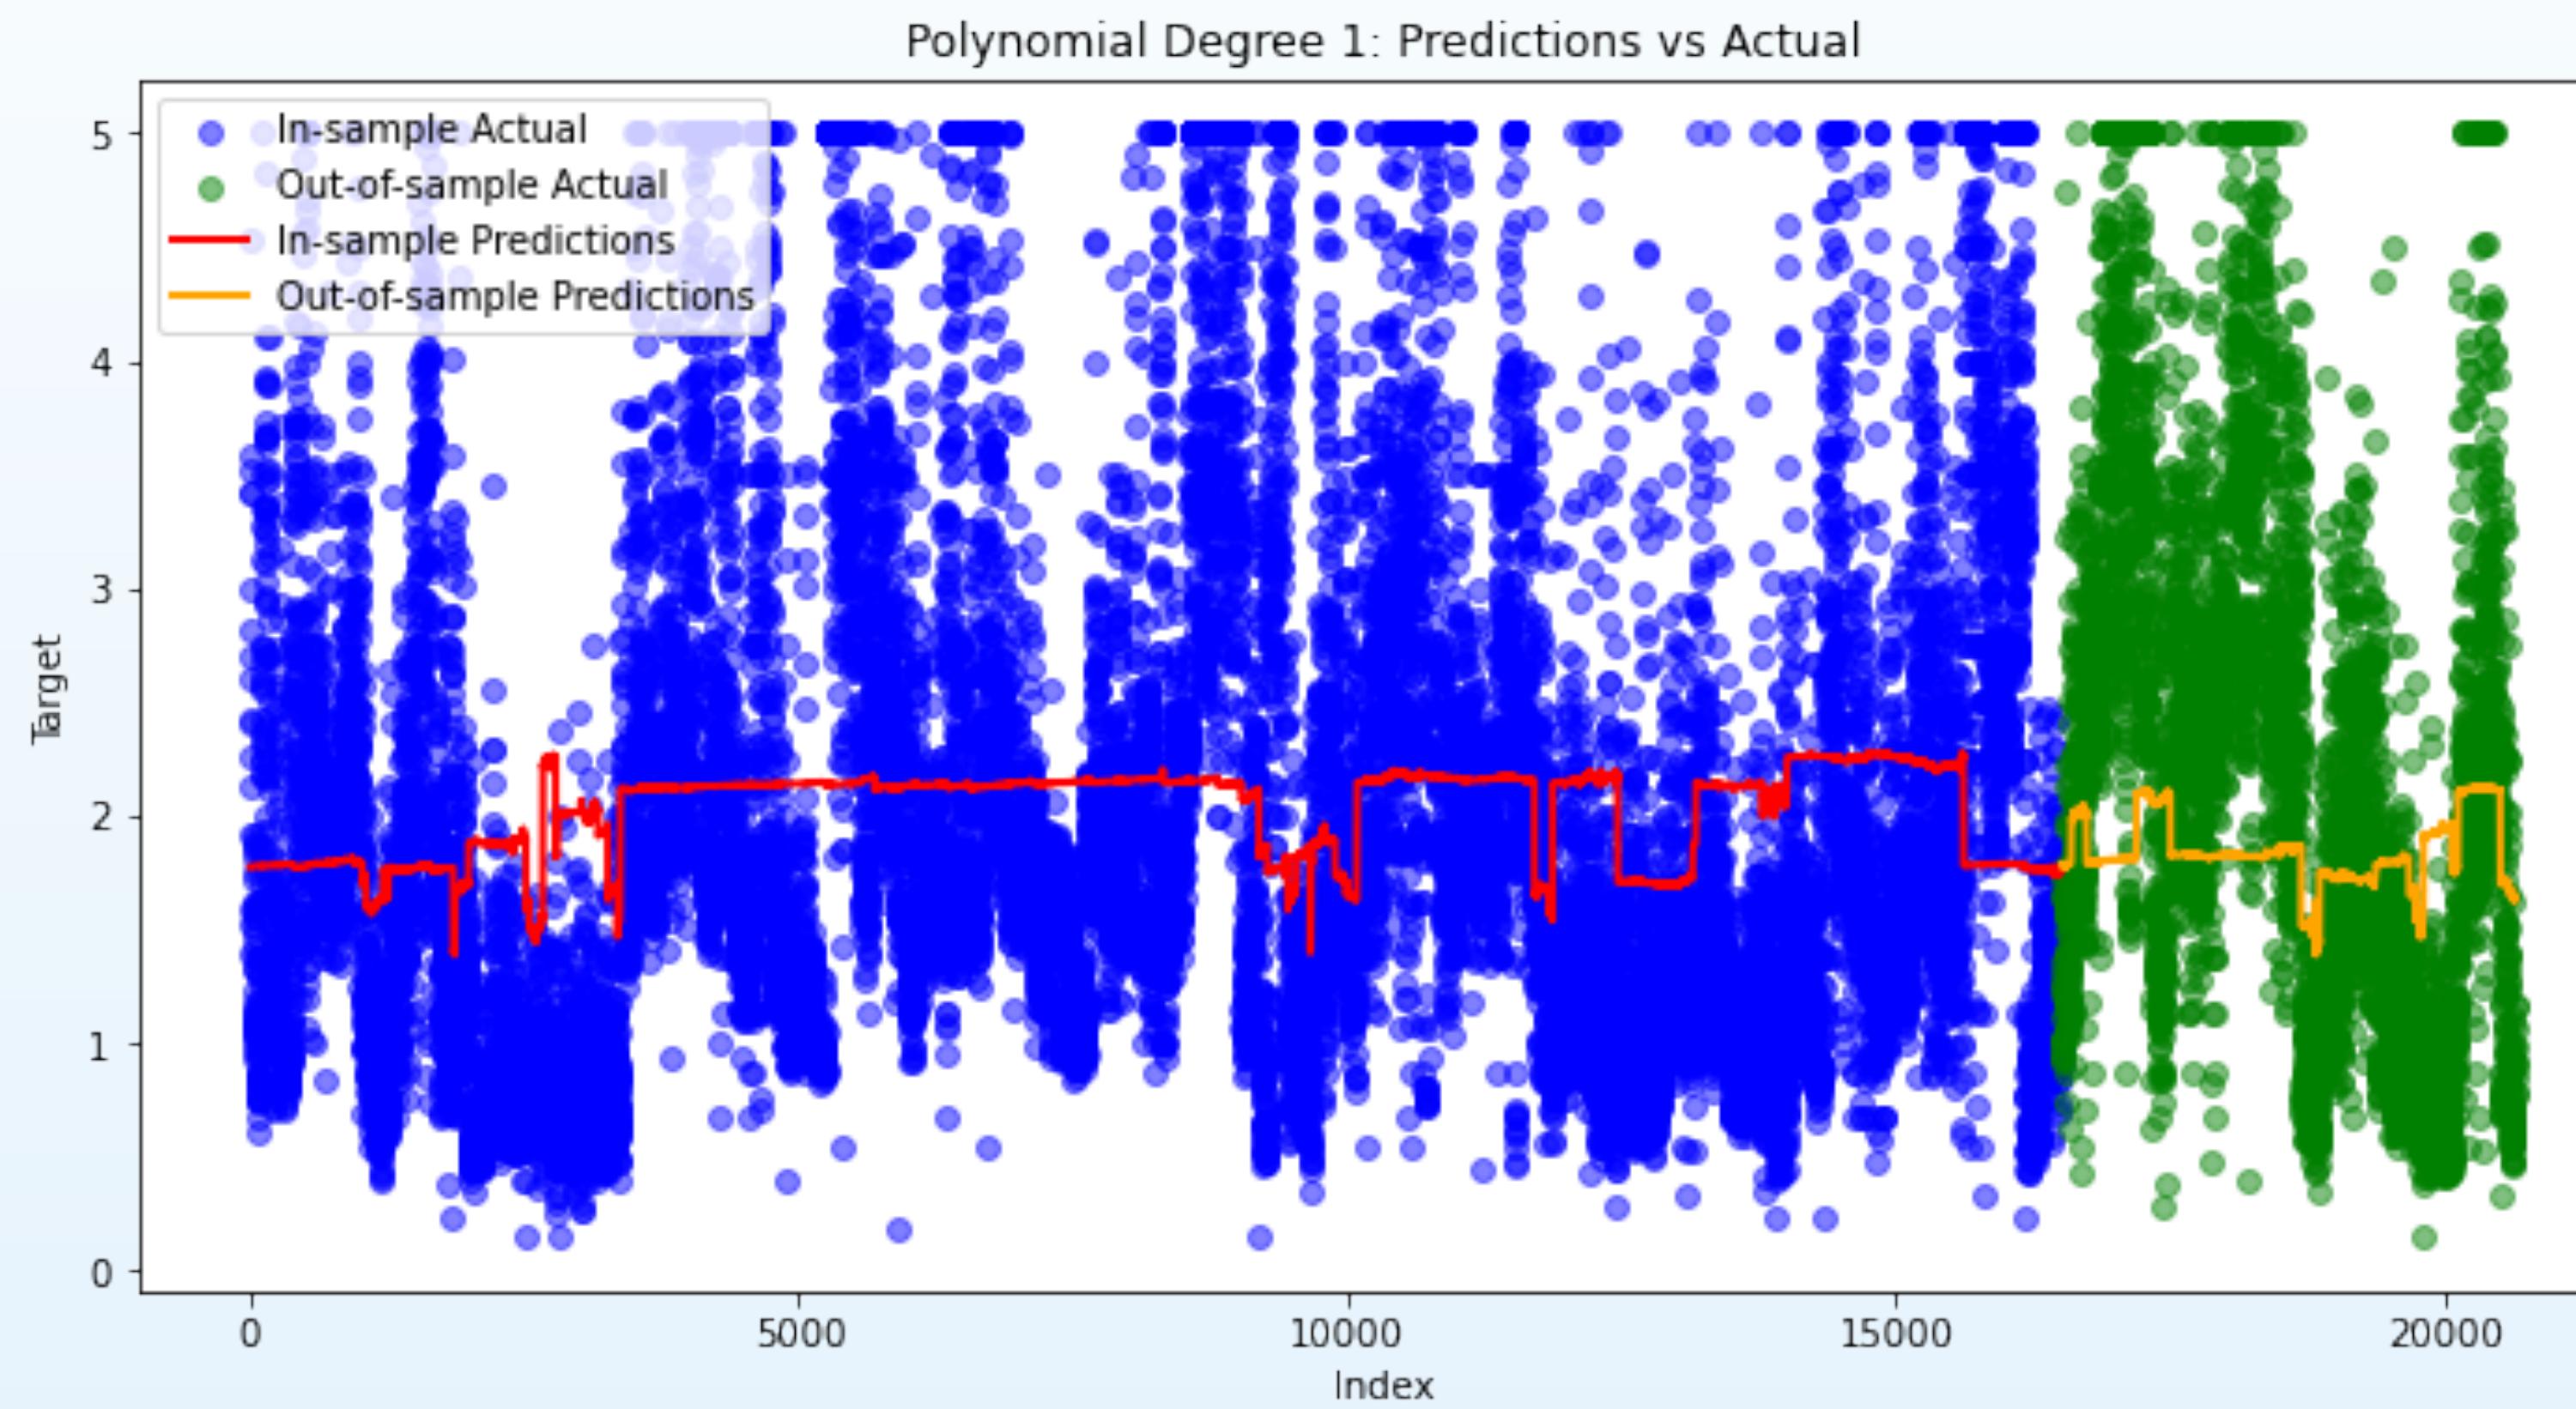
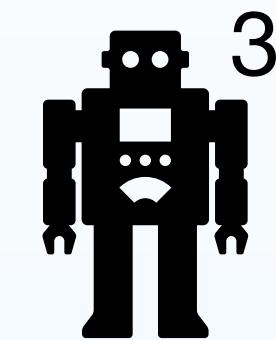
Various models are built using In-Sample Data with data best practices such as Cross Validation



# Machine Learning Pipeline : Example Regression



Final Model is selected  
based on Out-Of Sample  
performance

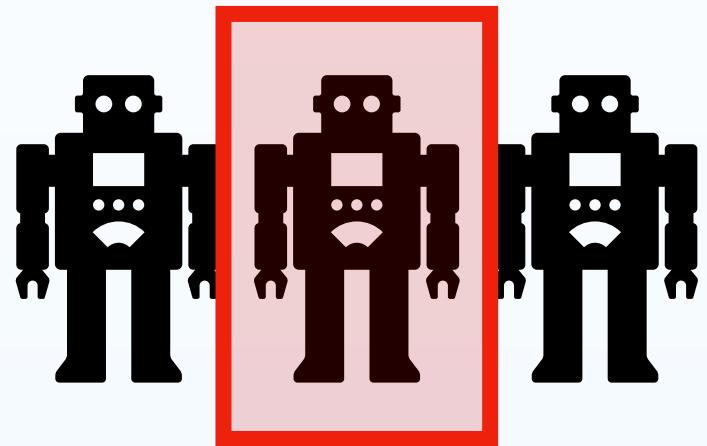


# Bias (Underfitting)

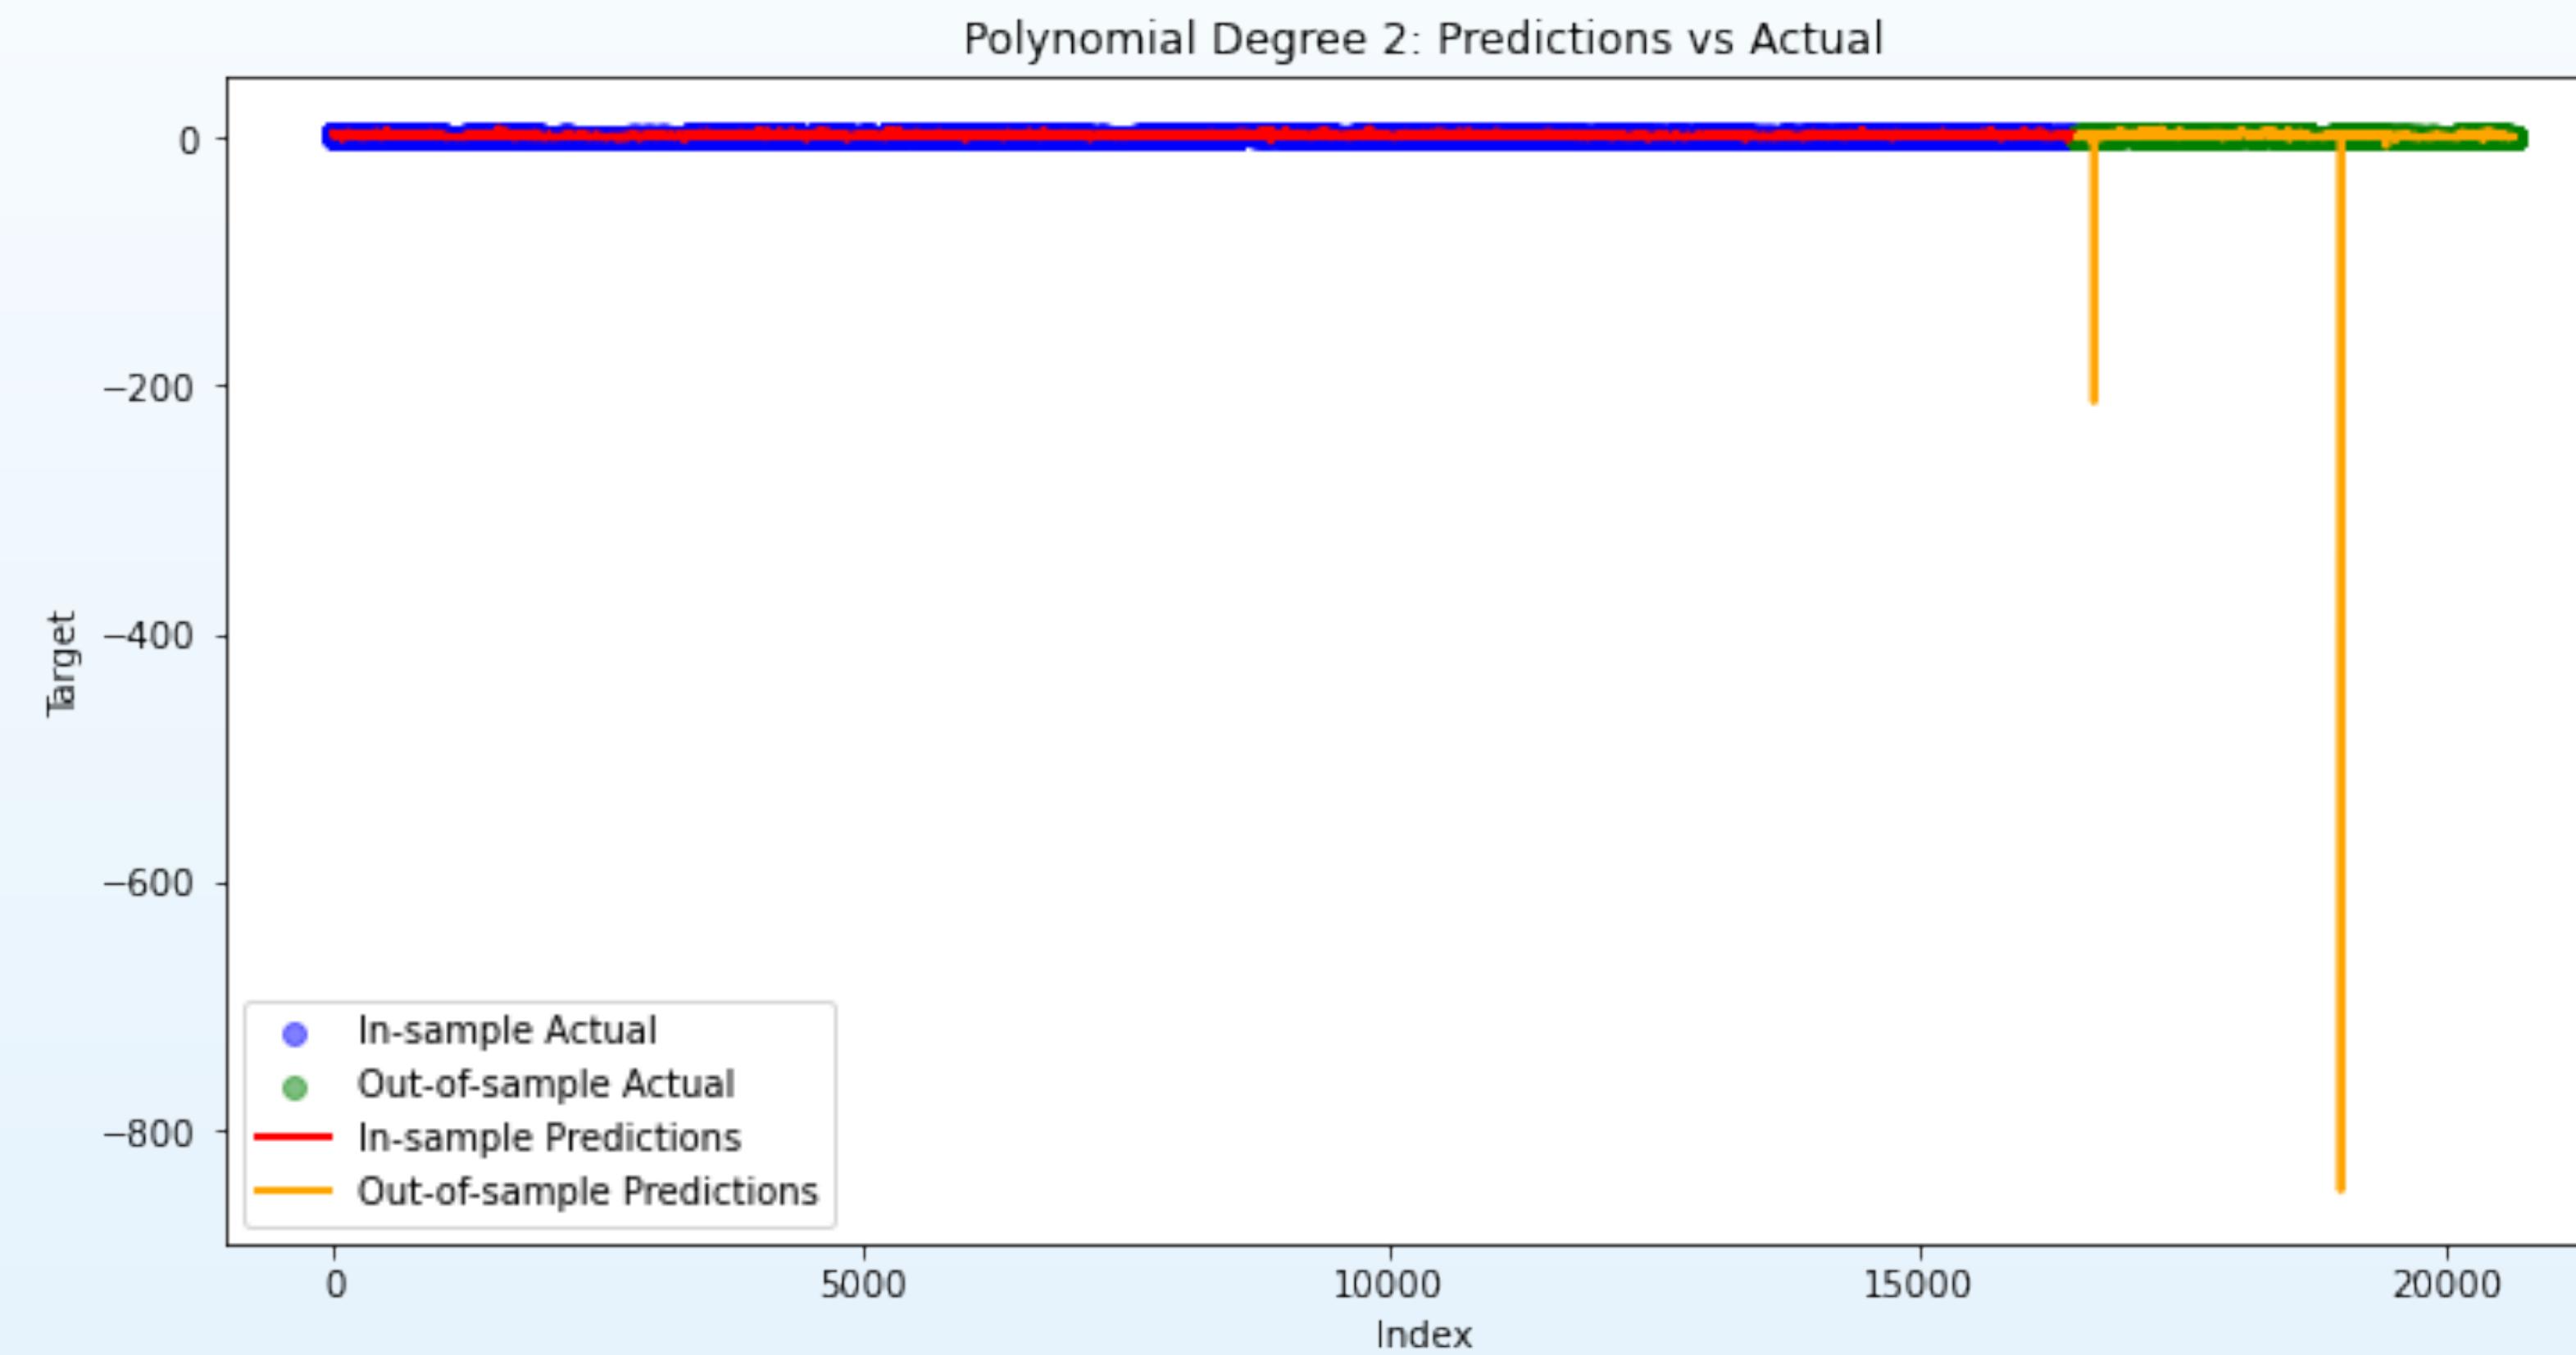
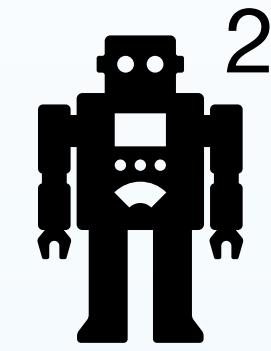
Bias refers to the error due to overly simplistic assumptions in the learning algorithm. A model with high bias tends to underfit the data, meaning it oversimplifies the underlying patterns. This leads to poor predictive performance as the model cannot capture the complexity of the real-world problem.

Example : Suppose you are training a linear regression model to predict housing prices. If the model assumes that the relationship between the features (e.g., square footage, number of bedrooms) and the price is strictly linear, it may perform poorly when the true relationship is more complex.

# Machine Learning Pipeline : Example Regression



Final Model is selected  
based on Out-Of Sample  
performance

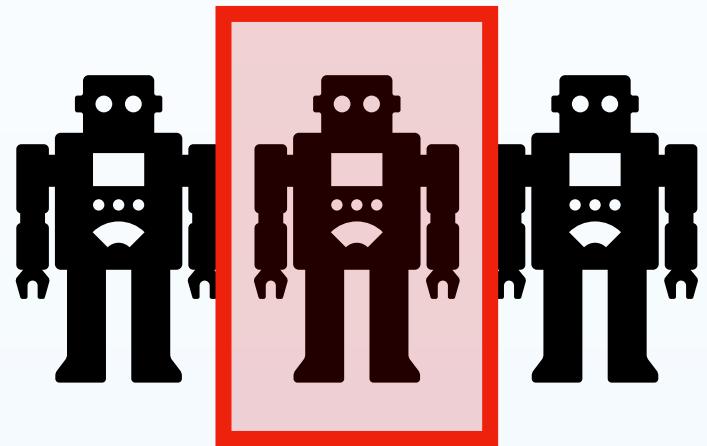


# Variance (Overfitting)

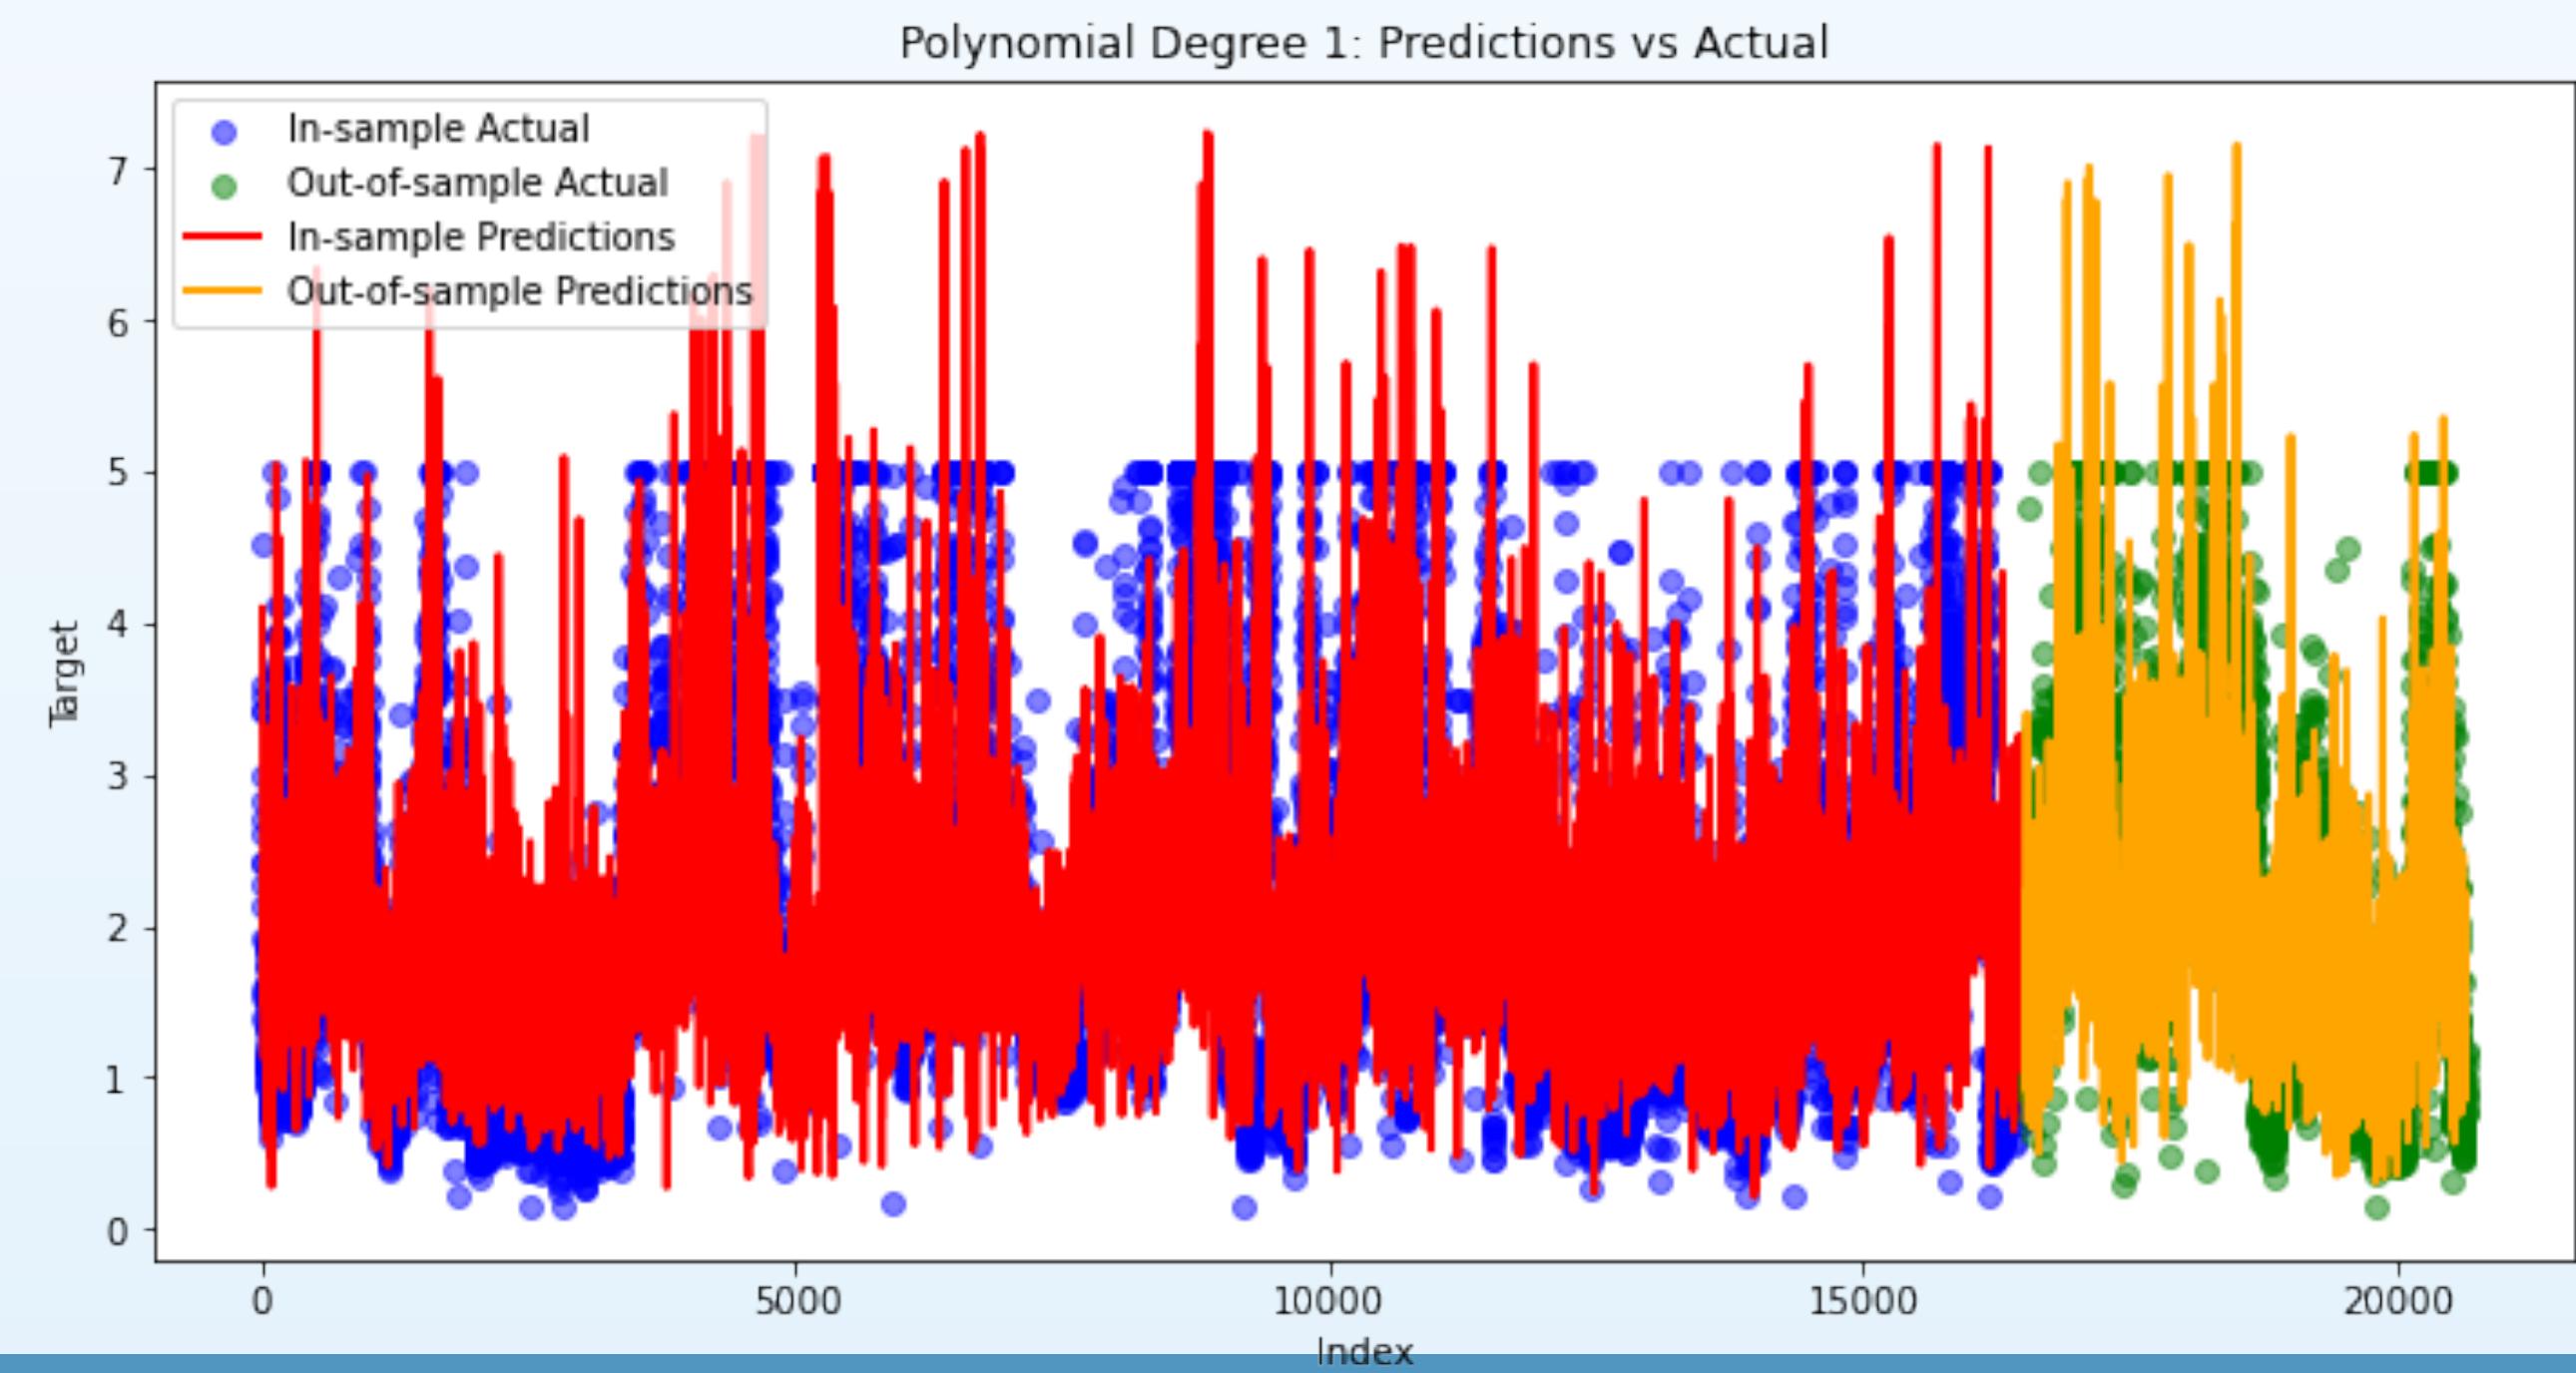
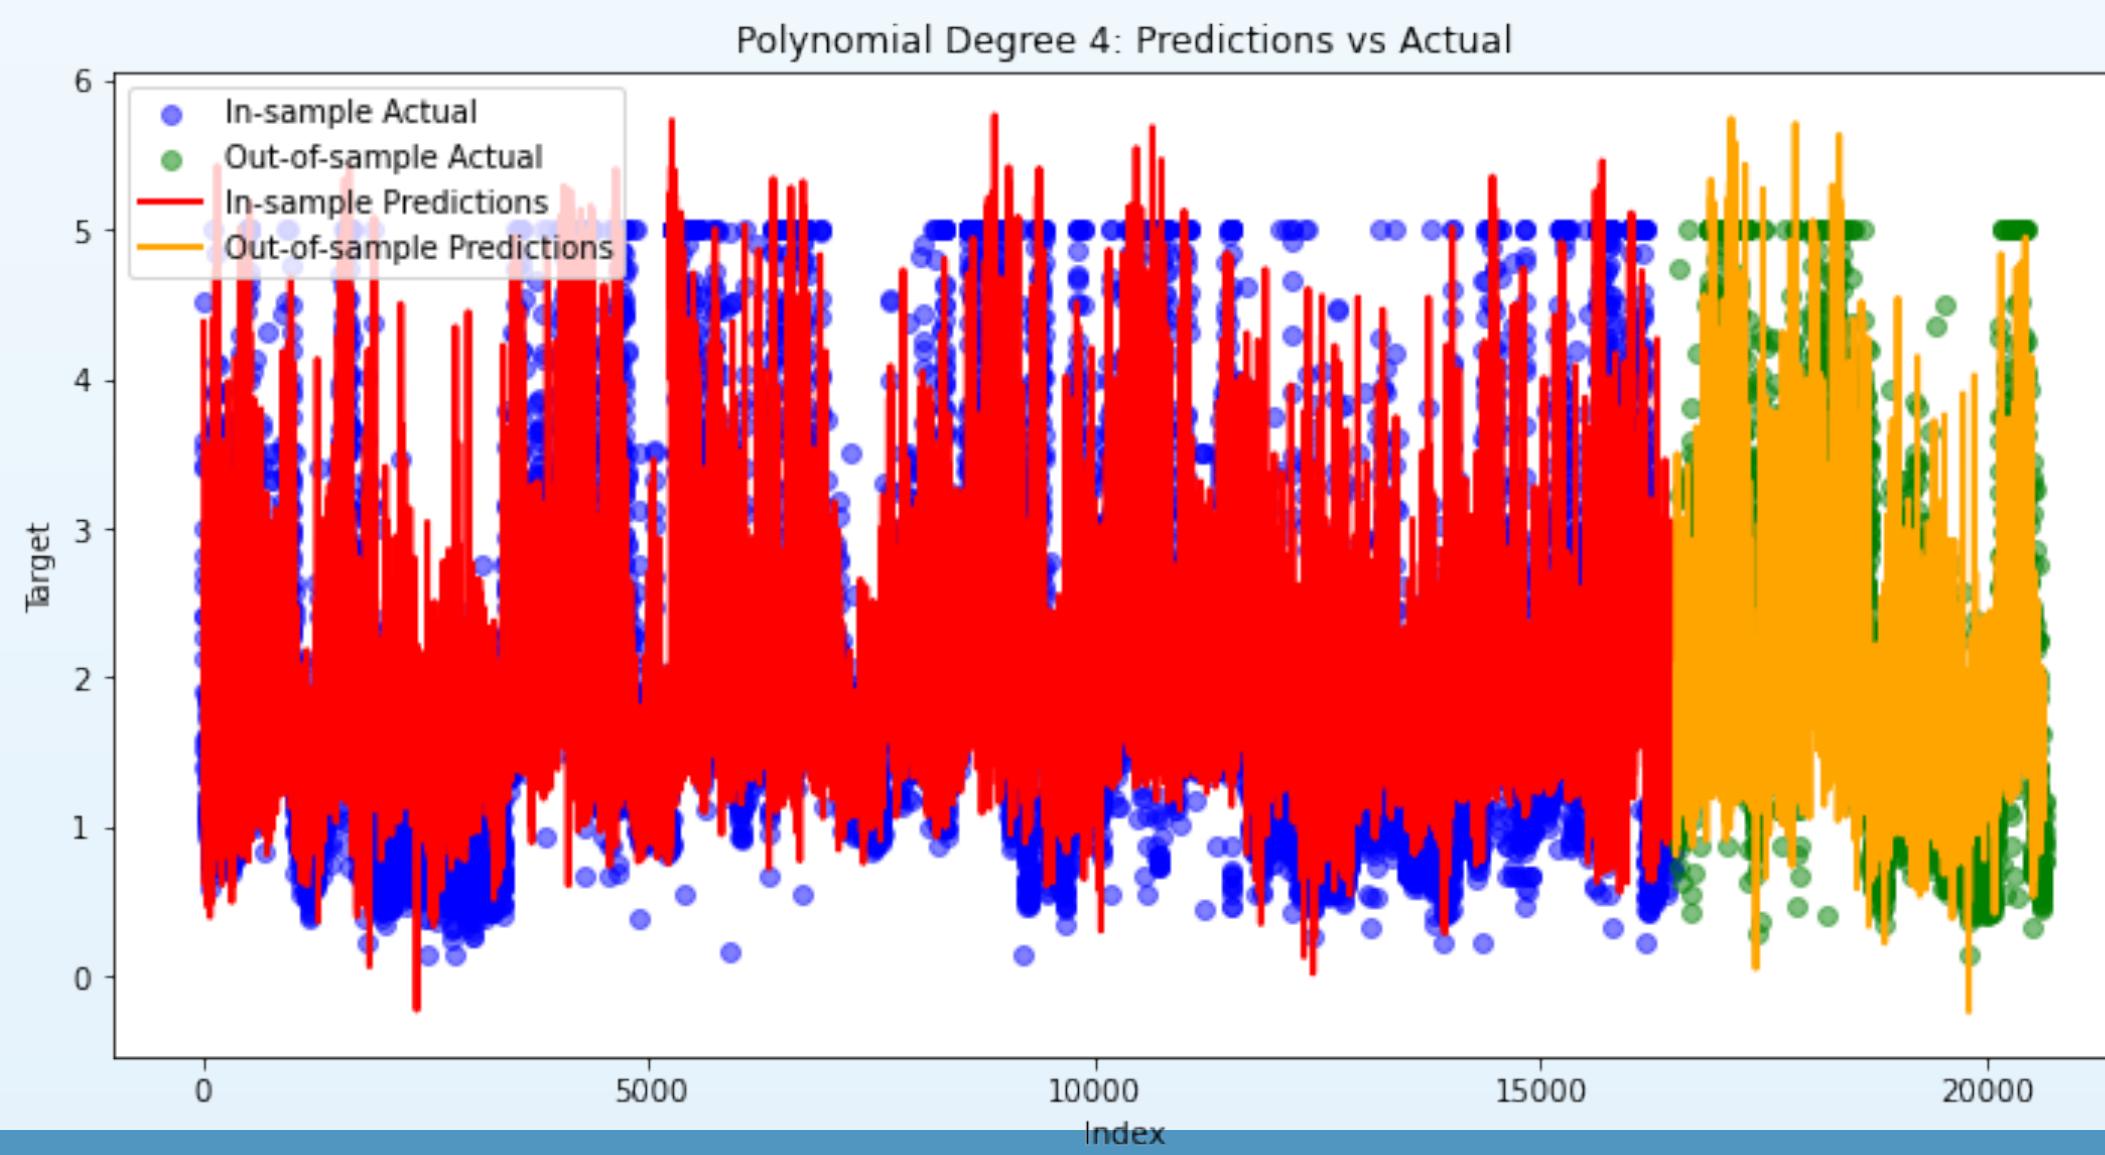
Variance, on the other hand, is the error due to excessive complexity in the learning algorithm. A model with high variance captures not only the underlying patterns but also the noise in the training data. This leads to poor generalization to unseen data.

Example : Imagine training a decision tree with a very deep structure on a dataset of handwritten digits. While the tree can fit the training data perfectly, it might perform poorly on new, unseen digits because it has essentially memorized the training examples, including their individual quirks.

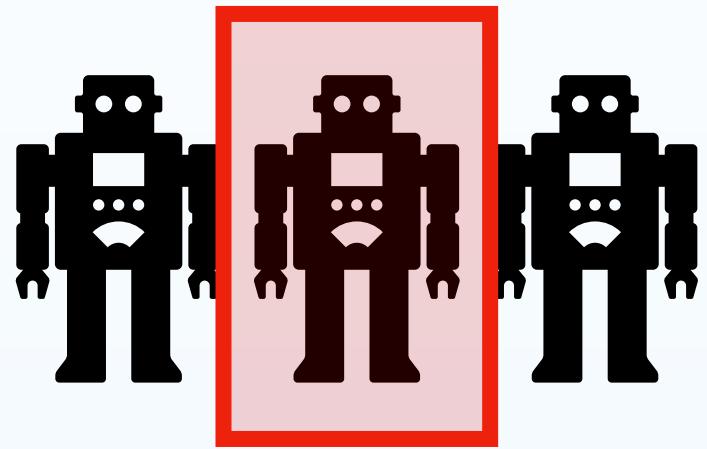
# Machine Learning Pipeline : Example Regression



Final Model is selected  
based on Out-Of Sample  
performance



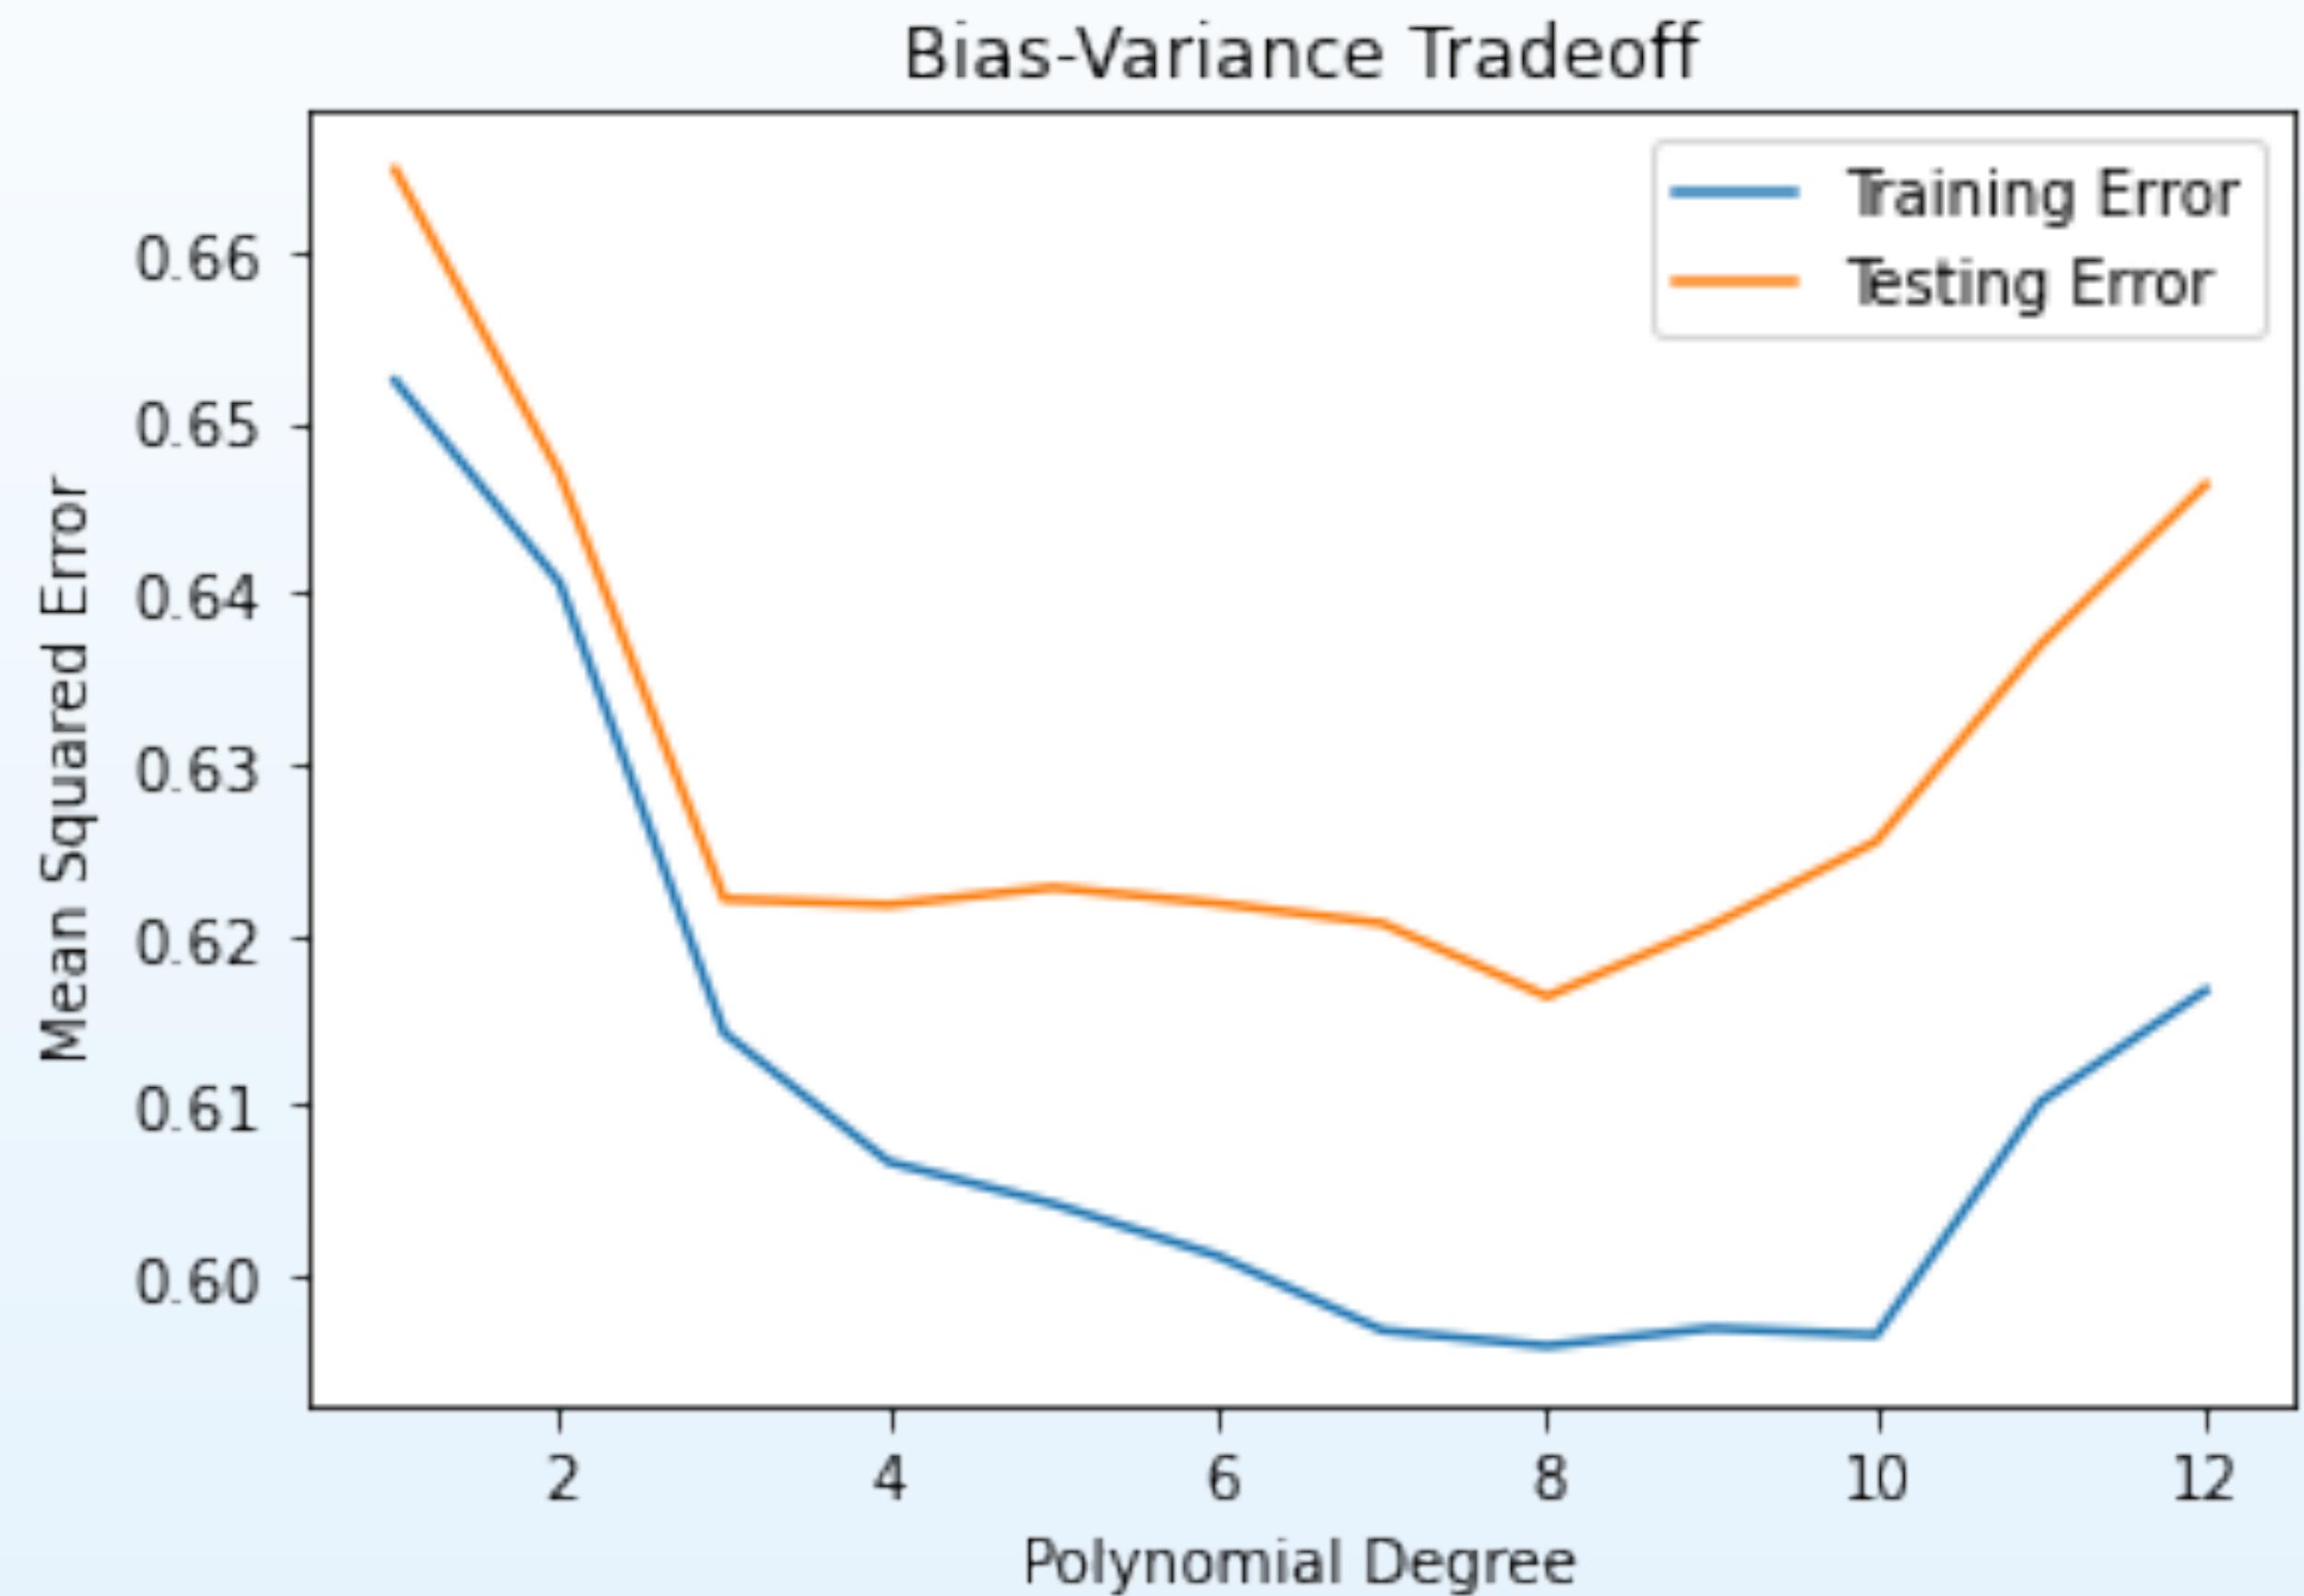
# Machine Learning Pipeline : Example Regression



Final Model is selected  
based on Out-Of Sample  
performance

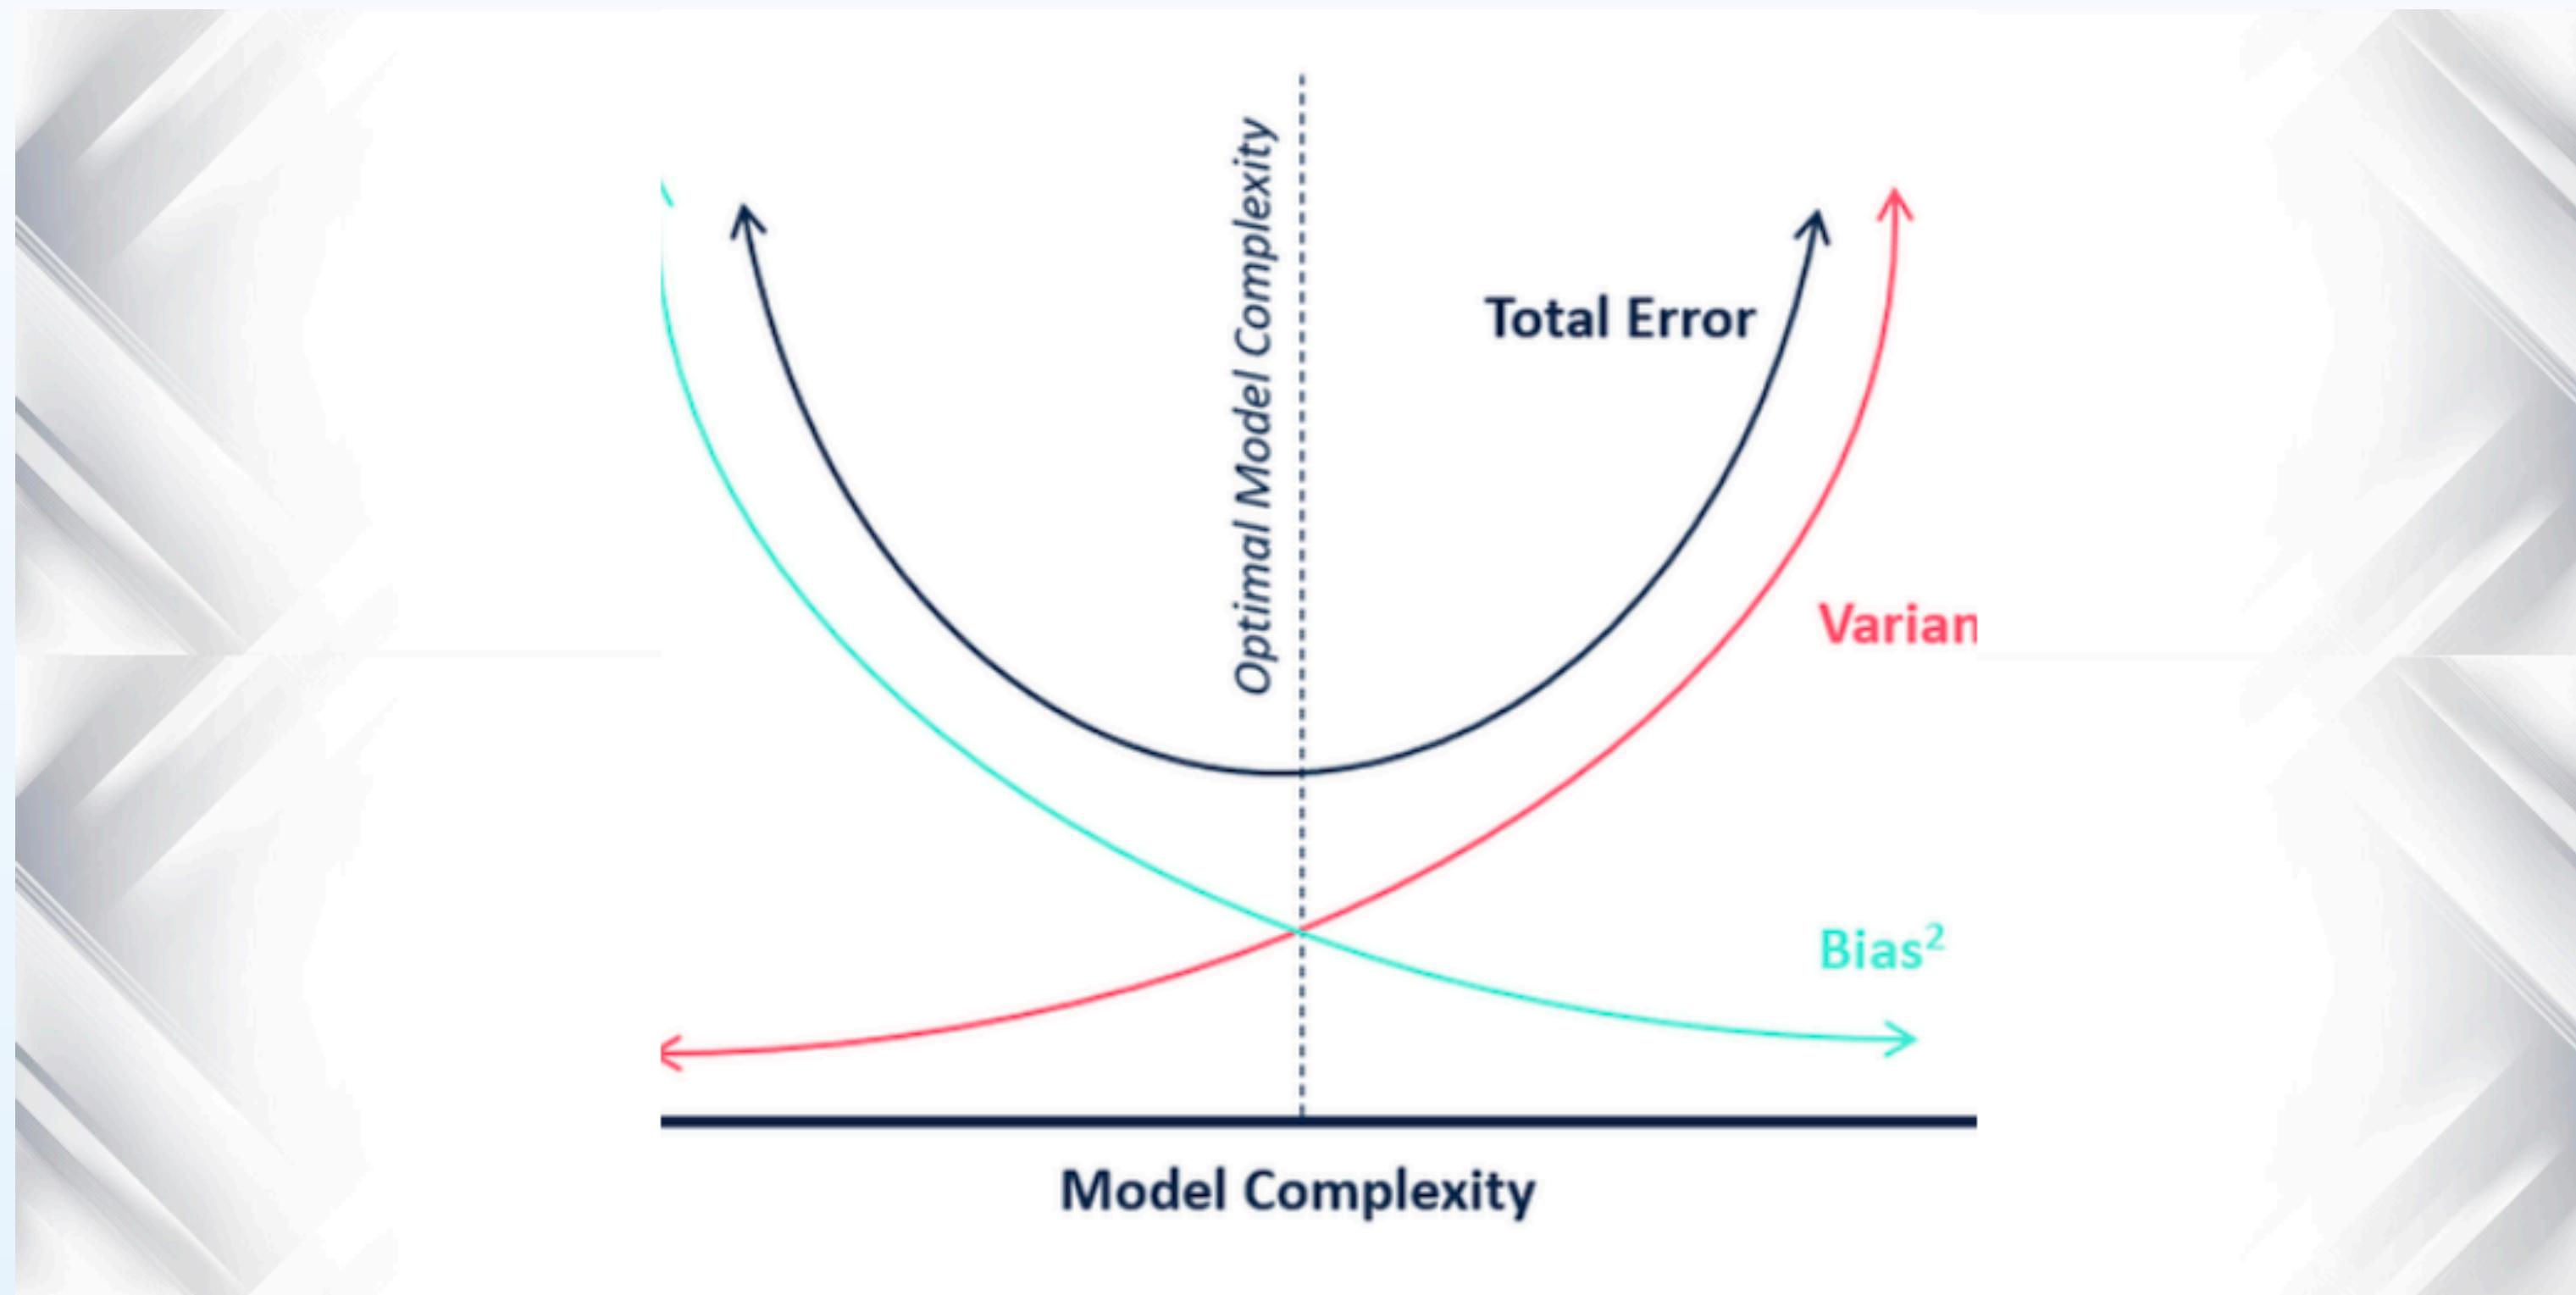
What is happening here?

As model complexity goes up we start getting higher Variance



# Measuring Bias Variance Trade-off

Always have to deal with the Bias-Variance trade off. We will learn models to optimize automatically



# Linear Regression with StatsModels

# Linear Regression

Have to meet the Gauss-Markov assumptions

We will not go through these in this class because you have done so much  
in 430

## The regression has five key assumptions:

- Linear relationship.
- Multivariate normality.
- No or little multicollinearity.
- No auto-correlation.
- Homoscedasticity.



# A python package that has prebuilt statistical test

Regression and Linear Models

Linear Regression

Module Reference

Generalized Linear Models

Generalized Estimating Equations

Generalized Additive Models (GAM)

Robust Linear Models

Linear Mixed Effects Models

Regression with Discrete Dependent Variable

Generalized Linear Mixed Effects Models

ANOVA

Other Models othermod

Time Series Analysis

Other Models

Statistics and Tools

Data Sets



We use scikit learn mainly in this class. So why stats models ever?

Mainly, their regression outputs

```
In [7]: print(res.summary())
OLS Regression Results
=====
Dep. Variable:          GRADE    R-squared:       0.416
Model:                 OLS     Adj. R-squared:   0.353
Method:                Least Squares F-statistic:    6.646
Date:      Thu, 04 Jan 2024 Prob (F-statistic): 0.00157
Time:          18:30:41   Log-Likelihood: -12.978
No. Observations:      32      AIC:             33.96
Df Residuals:          28      BIC:             39.82
Df Model:               3
Covariance Type:       nonrobust
=====
            coef    std err        t      P>|t|      [ 0.025    0.975 ]
-----
GPA         0.4639    0.162     2.864     0.008     0.132     0.796
TUCE        0.0105    0.019     0.539     0.594    -0.029     0.050
PSI         0.3786    0.139     2.720     0.011     0.093     0.664
const      -1.4980    0.524    -2.859     0.008    -2.571    -0.425
=====
Omnibus:           0.176   Durbin-Watson:    2.346
Prob(Omnibus):      0.916   Jarque-Bera (JB): 0.167
Skew:              0.141   Prob(JB):        0.920
Kurtosis:           2.786   Cond. No.       176.
=====
Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly
```

# Python Refresher

## Object

Anything that can be  
Stored as a variable  
Name

df

|   |      |
|---|------|
|   | 0    |
| 0 | 24.0 |
| 1 | 21.6 |
| 2 | 34.7 |

## Parameters

Creating an object or function,  
a user input

```
<func>( , <parameter> = <value>)  
  
pd.DataFrame(target,  
              columns = ["Target"])
```

## Methods

A function based  
For a specific object type  
<object>.<method>()

df.shift()

|   |      |
|---|------|
|   | 0    |
| 0 | NaN  |
| 1 | 24.0 |
| 2 | 21.6 |

## Attributes

An object associated  
with another object  
<object>.<attribute>

df.columns

RangeIndex(start=0,

# Splitting the Data

Pandas DataFrame

```
1 | d|
```

|       | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | Longitude | 0     |
|-------|--------|----------|----------|-----------|------------|----------|----------|-----------|-------|
| 0     | 8.3252 | 41.0     | 6.984127 | 1.023810  | 322.0      | 2.555556 | 37.88    | -122.23   | 4.526 |
| 1     | 8.3014 | 21.0     | 6.238137 | 0.971880  | 2401.0     | 2.109842 | 37.86    | -122.22   | 3.585 |
| 2     | 7.2574 | 52.0     | 8.288136 | 1.073446  | 496.0      | 2.802260 | 37.85    | -122.24   | 3.521 |
| 3     | 5.6431 | 52.0     | 5.817352 | 1.073059  | 558.0      | 2.547945 | 37.85    | -122.25   | 3.413 |
| 4     | 3.8462 | 52.0     | 6.281853 | 1.081081  | 565.0      | 2.181467 | 37.85    | -122.25   | 3.422 |
| ...   | ...    | ...      | ...      | ...       | ...        | ...      | ...      | ...       | ...   |
| 20635 | 1.5603 | 25.0     | 5.045455 | 1.133333  | 845.0      | 2.560606 | 39.48    | -121.09   | 0.781 |
| 20636 | 2.5568 | 18.0     | 6.114035 | 1.315789  | 356.0      | 3.122807 | 39.49    | -121.21   | 0.771 |
| 20637 | 1.7000 | 17.0     | 5.205543 | 1.120092  | 1007.0     | 2.325635 | 39.43    | -121.22   | 0.923 |
| 20638 | 1.8672 | 18.0     | 5.329513 | 1.171920  | 741.0      | 2.123209 | 39.43    | -121.32   | 0.847 |
| 20639 | 2.3886 | 16.0     | 5.254717 | 1.162264  | 1387.0     | 2.616981 | 39.37    | -121.24   | 0.894 |

20640 rows × 9 columns

Randomly Sampled

```
1 | d.sample(len(d))|
```

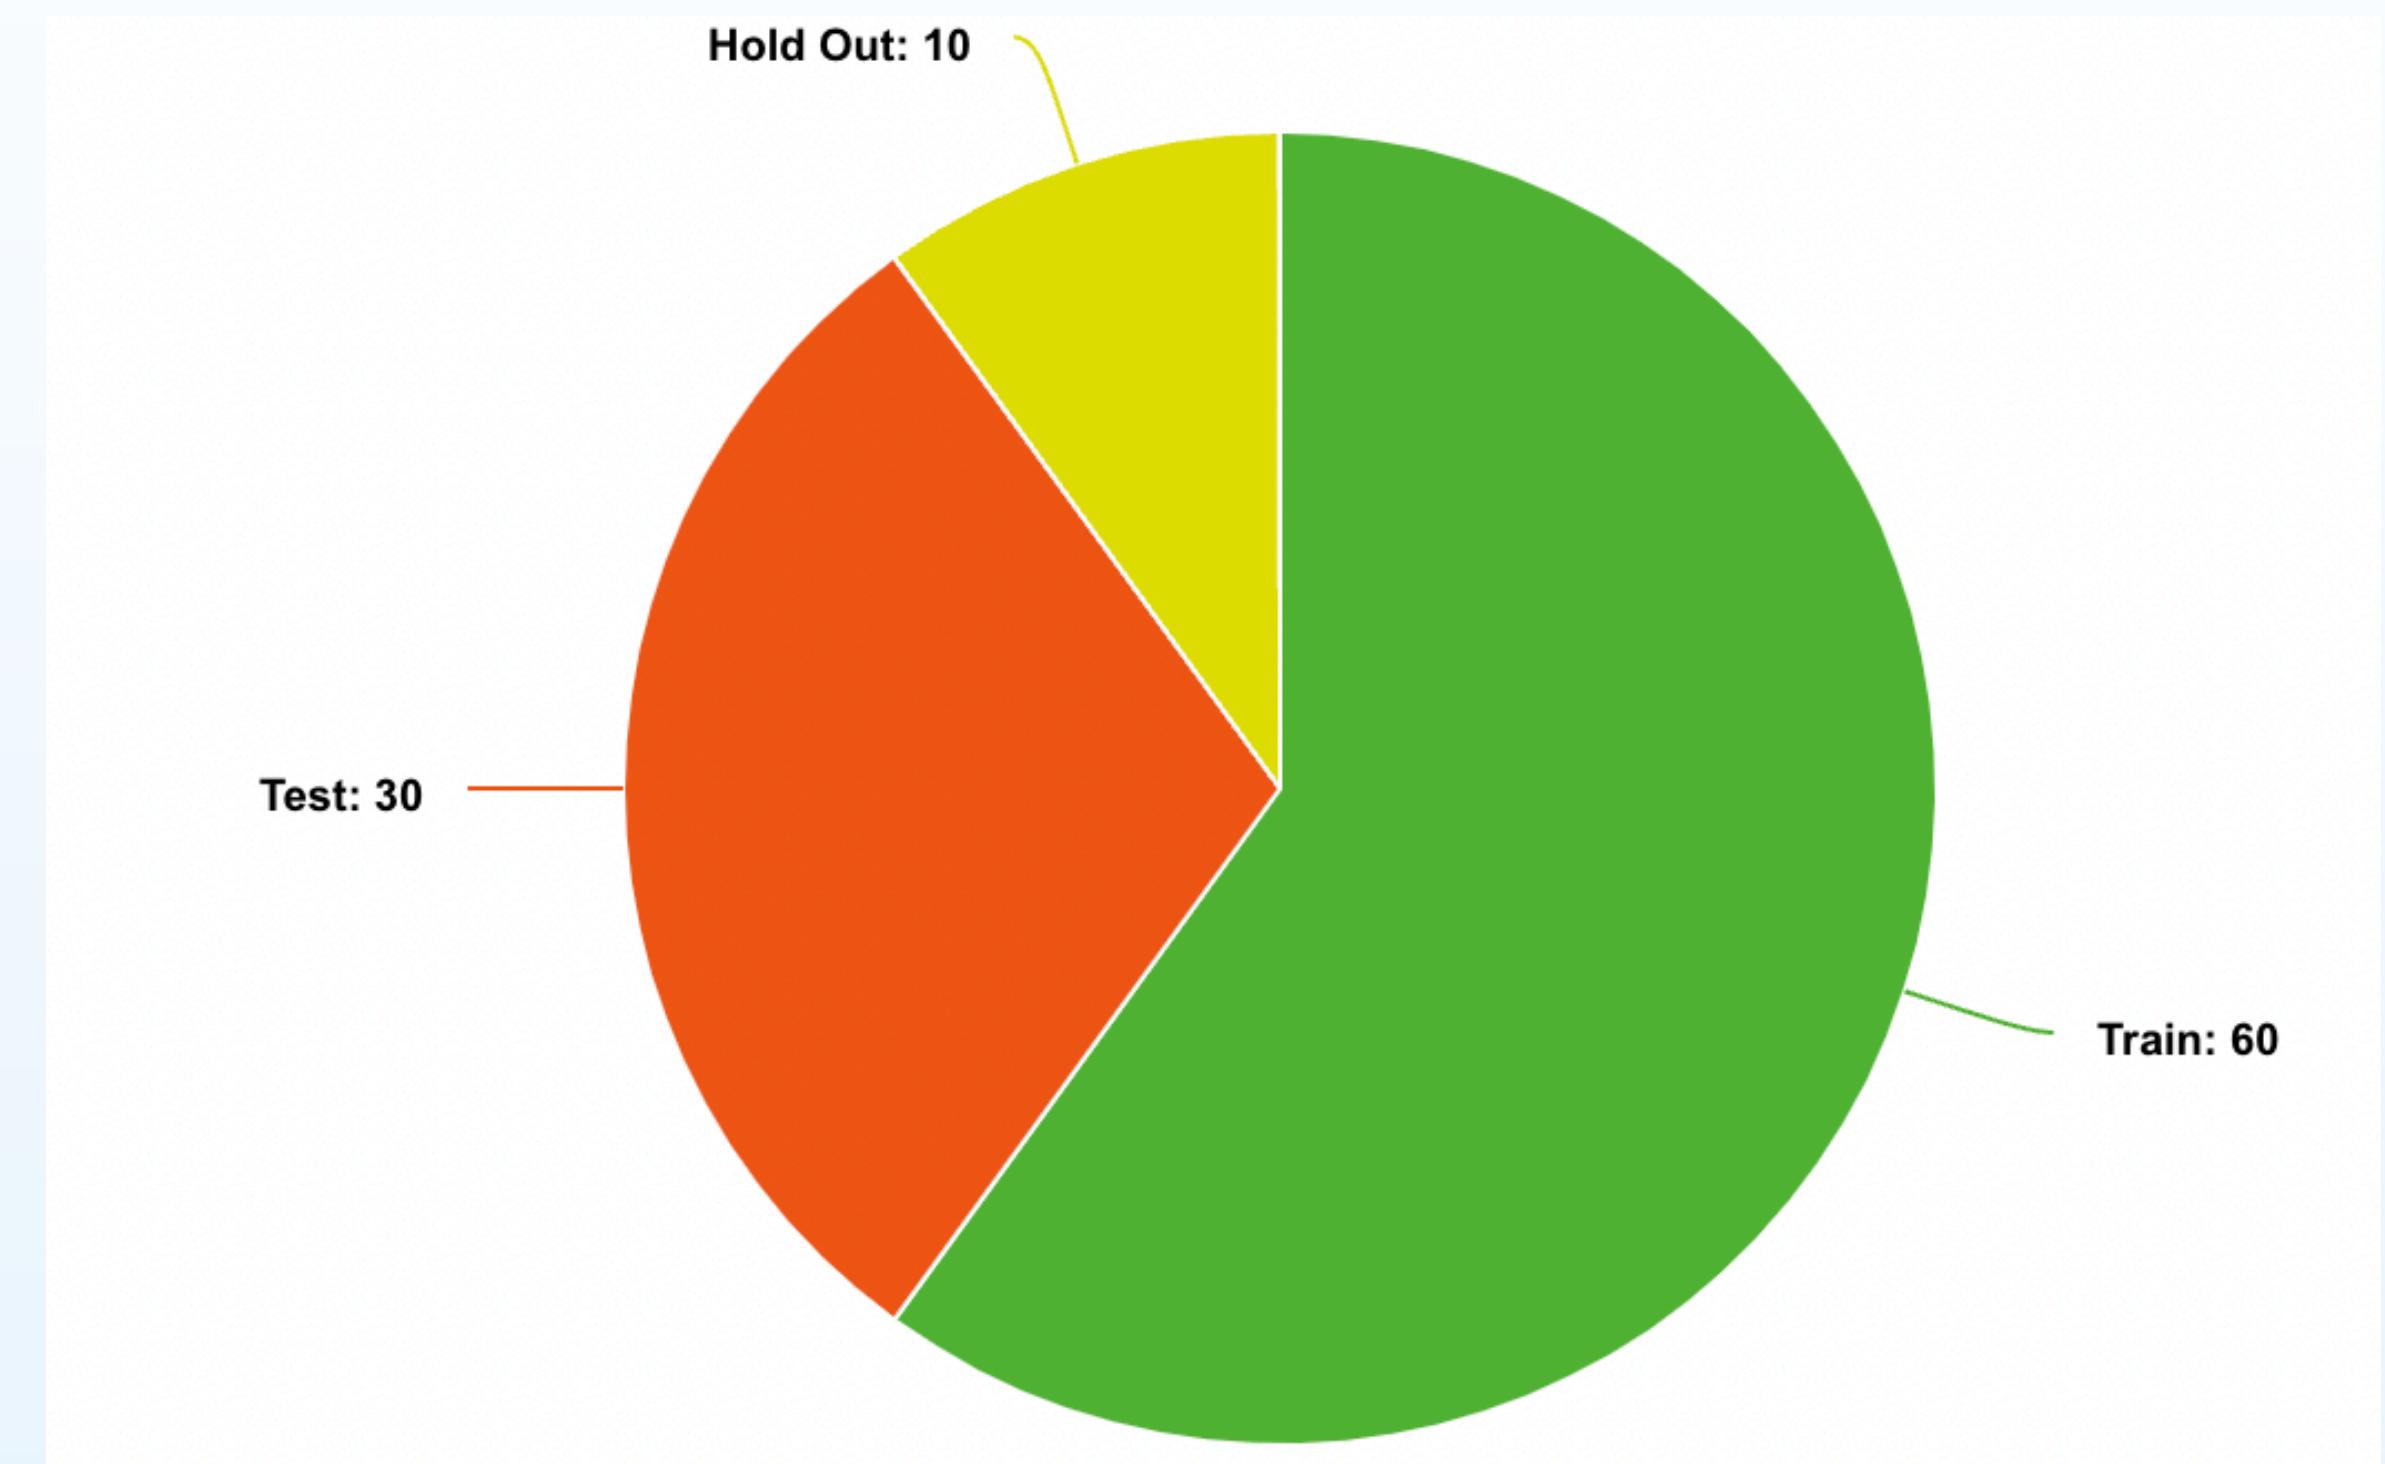
|       | MedInc  | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | Longitude | 0       |
|-------|---------|----------|----------|-----------|------------|----------|----------|-----------|---------|
| 6364  | 6.4576  | 33.0     | 5.661327 | 1.080092  | 1268.0     | 2.901602 | 34.15    | -117.97   | 5.00001 |
| 1181  | 2.2045  | 41.0     | 5.122222 | 1.188889  | 284.0      | 3.155556 | 39.48    | -121.55   | 0.41800 |
| 13742 | 10.8634 | 20.0     | 5.903448 | 0.855172  | 395.0      | 2.724138 | 34.03    | -117.19   | 3.81800 |
| 629   | 3.3710  | 42.0     | 4.530120 | 0.992470  | 1817.0     | 2.736446 | 37.72    | -122.17   | 1.65000 |
| 3858  | 11.2093 | 42.0     | 6.876106 | 0.902655  | 284.0      | 2.513274 | 34.17    | -118.43   | 5.00001 |
| ...   | ...     | ...      | ...      | ...       | ...        | ...      | ...      | ...       | ...     |
| 10345 | 5.7871  | 24.0     | 6.184295 | 1.059295  | 1972.0     | 3.160256 | 33.61    | -117.67   | 2.27400 |
| 5001  | 1.7849  | 46.0     | 4.294326 | 1.138298  | 1153.0     | 4.088652 | 33.99    | -118.28   | 0.99300 |
| 15172 | 5.7382  | 17.0     | 6.444043 | 0.920578  | 667.0      | 2.407942 | 33.04    | -117.06   | 2.78000 |
| 11318 | 3.4955  | 15.0     | 3.709447 | 0.946524  | 1617.0     | 2.882353 | 33.77    | -117.99   | 1.60900 |
| 19869 | 3.1486  | 14.0     | 5.218341 | 0.960699  | 568.0      | 2.480349 | 36.33    | -119.35   | 1.05600 |

20640 rows × 9 columns

# Splitting the Data

```
split1 = int(len(df) * .6)
split2 = int(len(df) * .9)
d_in = d[:split1]
d_out = d[split1:split2]
d_hold = d[split2:]
```

```
X_in = d_in.iloc[:, :-1]
y_in = d_in.iloc[:, -1:]
```



# Running the Model

Build your first Model

## INPUT TESTING OF GUASS-MARKOV ASSUMPTIONS

```
X_in = sm.add_constant(X_in)
model1 = sm.OLS(y_in, X_in).fit()
```

```
dir(model1)
```

```
'aic',
'bic',
'bse',
'centered_tss',
'compare_f_test',
'compare_lm_test',
'compare_lr_test',
'condition_number',
'conf_int',
'conf_int_el',
'cov_HC0',
'cov_HC1',
'cov_HC2',
'cov_HC3',
'cov_kwds',
'cov_params',
'cov_type',
'df_model',
'df_resid',
'eigenvals',
'el_test',
'ess',
'f_pvalue',
'f_test',
'fittedvalues',
'fvalue',
'get_influence',
'get_prediction',
'get_robustcov_results',
'info_criteria',
'initialize',
'k_constant',
'llf',
'load',
'model',
'mse_model',
'mse_resid',
'mse_total',
'nobs',
'normalized_cov_params',
'outlier_test',
'params',
'predict',
'pvalues',
'remove_data',
'resid',
'resid_pearson',
'rsquared',
'rsquared_adj',
'save',
'scale',
'ssr',
'summary',
'summary2',
't_test',
't_test_pairwise',
'tvalues',
'uncentered_tss',
'use_t',
'wald_test',
'wald_test_terms',
'wresid']
```

# Analyzing the Model

## Build your first Model

```
1 print(model1.summary())
```

| OLS Regression Results |                  |                     |           |          |           |           |
|------------------------|------------------|---------------------|-----------|----------|-----------|-----------|
| Dep. Variable:         | 0                | R-squared:          | 0.589     |          |           |           |
| Model:                 | OLS              | Adj. R-squared:     | 0.589     |          |           |           |
| Method:                | Least Squares    | F-statistic:        | 2216.     |          |           |           |
| Date:                  | Wed, 10 Jan 2024 | Prob (F-statistic): | 0.00      |          |           |           |
| Time:                  | 11:31:26         | Log-Likelihood:     | -13748.   |          |           |           |
| No. Observations:      | 12384            | AIC:                | 2.751e+04 |          |           |           |
| Df Residuals:          | 12375            | BIC:                | 2.758e+04 |          |           |           |
| Df Model:              | 8                |                     |           |          |           |           |
| Covariance Type:       | nonrobust        |                     |           |          |           |           |
| coef                   | std err          | t                   | P> t      | [0.025   | 0.975]    |           |
| const                  | -28.7672         | 1.004               | -28.641   | 0.000    | -30.736   | -26.798   |
| MedInc                 | 0.4330           | 0.005               | 80.271    | 0.000    | 0.422     | 0.444     |
| HouseAge               | 0.0056           | 0.001               | 9.020     | 0.000    | 0.004     | 0.007     |
| AveRooms               | -0.0949          | 0.008               | -12.531   | 0.000    | -0.110    | -0.080    |
| AveBedrms              | 0.5781           | 0.036               | 16.135    | 0.000    | 0.508     | 0.648     |
| Population             | -1.546e-05       | 6.34e-06            | -2.437    | 0.015    | -2.79e-05 | -3.02e-06 |
| AveOccup               | -0.0080          | 0.001               | -6.659    | 0.000    | -0.010    | -0.006    |
| Latitude               | -0.4003          | 0.011               | -36.586   | 0.000    | -0.422    | -0.379    |
| Longitude              | -0.3611          | 0.011               | -31.449   | 0.000    | -0.384    | -0.339    |
| Omnibus:               | 2751.880         | Durbin-Watson:      |           | 0.842    |           |           |
| Prob(Omnibus):         | 0.000            | Jarque-Bera (JB):   |           | 8334.861 |           |           |
| Skew:                  | 1.147            | Prob(JB):           |           | 0.00     |           |           |
| Kurtosis:              | 6.301            | Cond. No.           |           | 2.76e+05 |           |           |

```
1 model1.rsquared_adj
```

0.5886974211054199

```
1 model1.mse_total
```

1.3119119613283674

```
1 model1.params
```

|            |            |
|------------|------------|
| const      | -28.767223 |
| MedInc     | 0.432963   |
| HouseAge   | 0.005554   |
| AveRooms   | -0.094948  |
| AveBedrms  | 0.578070   |
| Population | -0.000015  |
| AveOccup   | -0.007995  |
| Latitude   | -0.400267  |
| Longitude  | -0.361127  |

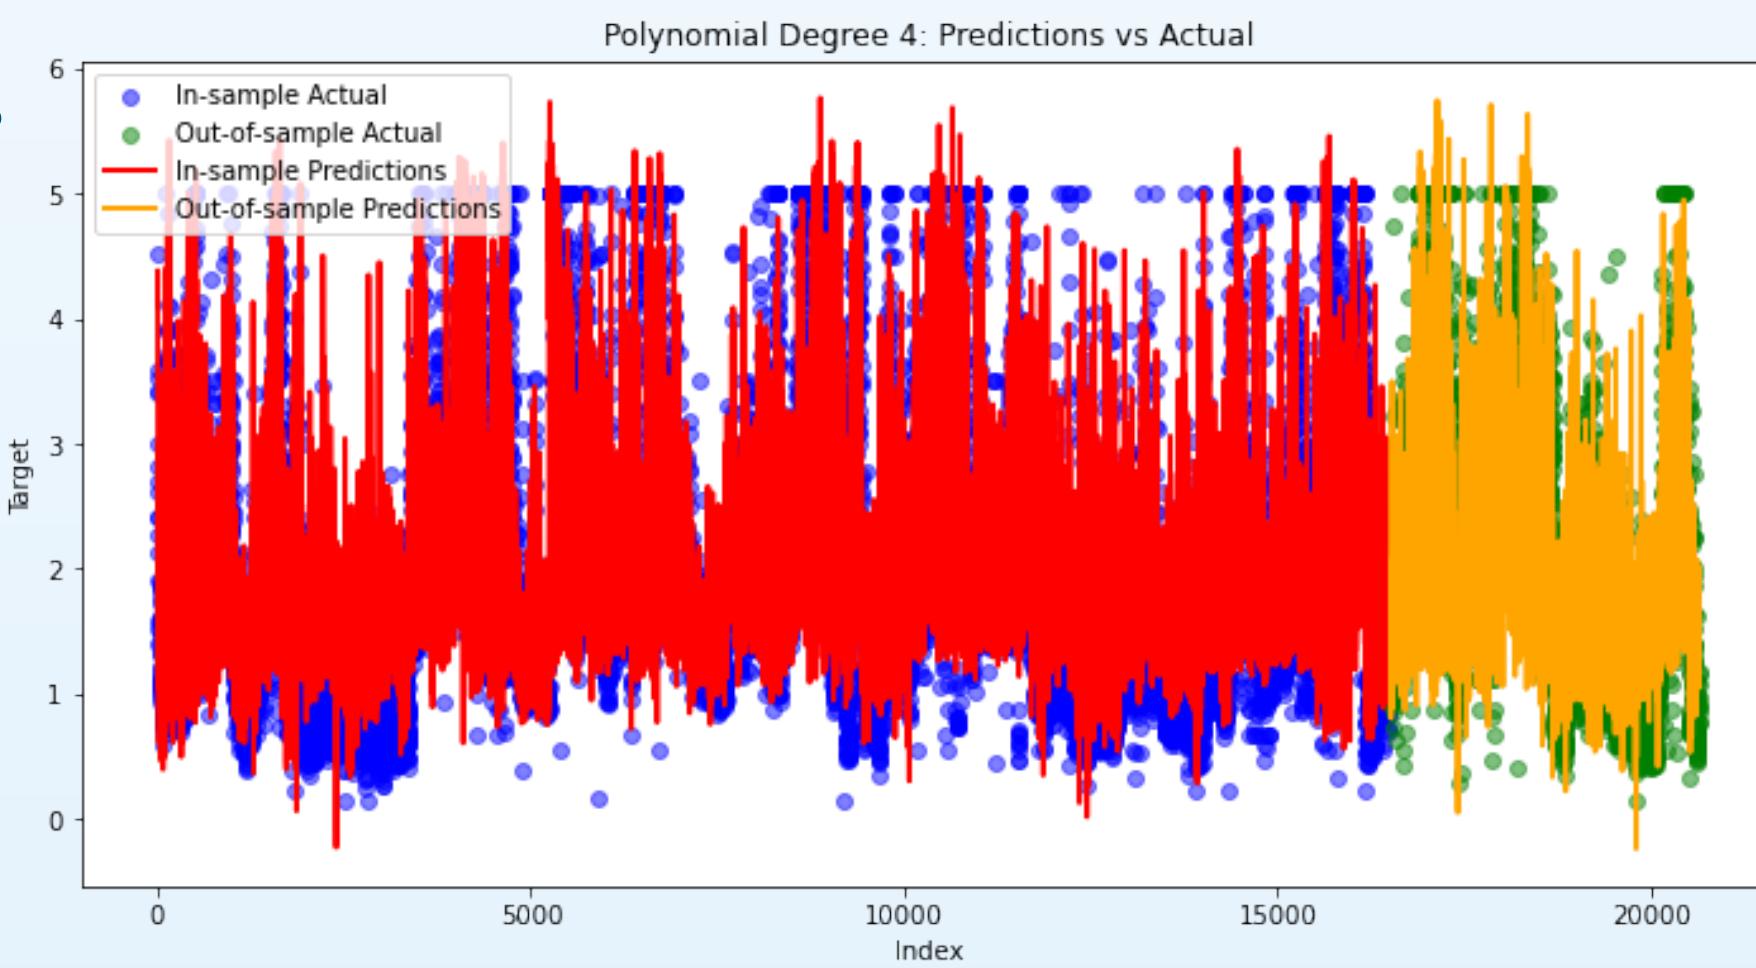
# Test in Out of Sample

Pick your measure. Which ever does the best in out of sample is your selected model

```
out_pred = model1.predict(X_out)  
mean_squared_error(y_out, out_pred)
```

# In-Class Assignment

- o.) Make a Jupyter notebook
- i.) Import from FRED Data
- 2.) Do not randomize, split the data into Train test hold out
- 3.) Build a model that regresses FF-Unemp, HousingStarts, Inflation
- 4.) Recreate this graph for your model
- 5.) What is the in and out of Sample MSE.



- 6.) Using a for loop. Repeat 3,4,5 for polynomial degrees 1,2,3
- 7.) State your observations